

# LAB : OpenMP



--

Texas Advanced Computing Center  
The University of Texas at Austin

2016/03/31

Sides at: [tinyurl.com/tacc-openmp](http://tinyurl.com/tacc-openmp) or [portal.tacc.utexas.edu/training](http://portal.tacc.utexas.edu/training) click View Details

# Introduction

## What you will learn

- How to compile Code (C and Fortran) with OpenMP
- How to parallelize code with OpenMP
  - Use the correct header declarations
  - Parallelize simple loops
- How to effectively hide OpenMP statements

## What you will do

- Modify example code **READ the CODE COMMENTS**
- Compile and execute the example
- Compare the run-time of the serial codes and the OpenMP parallel codes with different scheduling methods

# Accessing Lab Files

- Log on to **Lonestar 5** using your account.
- Untar the file lab\_Openmp.tar file (in ~train00).
- The new directory (lab\_openmp) contains sub-directories for exercises .
- cd into the appropriate subdirectory for an exercise.

```
ssh <your.login.name>@ls5.tacc.utexas.edu  
tar -xvf ~train00/lab_OpenMP.tar  
cd lab_openmp
```

# Running on compute nodes Interactively

You can do the labs without submitting batch jobs

- You can compile\* and execute your code on the login node (login#.ls5); or you can use one of the compute nodes (c###-###). Here is how to do that.

```
1.          60 minutes    account
login2$ iddev -m 60 -A TRAINING-HPC
--> Verifying availability of home(/home1/00770/...) ...OK
--> Verifying access to desired queue (devel)          ...OK
c559-802$
```

2. Once you have a command prompt, you are ready to go (you own the node- it isn't shared with any other user).  
E.g. compile and execute - note the login prompt is the node name.

```
(this is only an example)
c559-001% ifort hello.f90 -o hello
c559-001% ./hello
```

# Compiling and running

- All OpenMP statements are activated by the OpenMP flag:
  - Intel compiler: `icc/ifort -fopenmp -fpp source.<c,f90>`
- **We will be using the Intel compiler**
- Compilation with the OpenMP flag (`-openmp`):
  - Activates OpenMP comment directives (...):
    - Fortran: `!$OMP ...`
    - C: `#pragma omp ...`
  - Enables the macro named `_OPENMP`
    - `#ifdef _OPENMP` evaluates to true  
(Fortraners: compile with `-fpp`)
  - Enables "hidden" statements (Fortran only!)
    - `!$ ...`
- Control runs with `OMP_NUM_THREADS`

# Exercise: daxpy

- Easy to parallelize loop
- Use directives
- Measure speedup

# Exercise: daxpy

- cd exercise\_daxpy
- Codes: f\_daxpy.f90/c\_daxpy.c
- Number of intervals is varied (Trial loop)

daxpy

**Trial Loop:    *itr*ial**  
**Loop over *i***

1 Parallelize the Loop over *i* :  
Use **omp parallel do/for**  
with the default(none) clause

2 Compile with:  
**make f\_daxpy**  
or  
**make c\_daxpy**

3 Run with 1 and 24

4 Compare timings

- Why is performance only doubled?

- Parallelize the code

- 1 complete OpenMP statements
- Initialization
  - **omp get max threads**

- ✓ Hint: Parallel performance can be limited by memory bandwidth– what is happening for every daxpy operation? (Is there cache reuse?)

# Exercise: pi

- Calculation of pi by numerical integration:  
lots of small steps
- Use parallel do
- Problem how to deal with sum reduction
- Wrong approach: critical section
- Right approach: use reduction clause



# Exercise : $\pi$ Integration

- cd exercise\_pi
- Codes: f\_pi.F90/c\_pi.c
- Number of intervals is varied (Trial loop)

## $\pi$ calculation

**Trial Loop:  $i_{\text{trial}}$**   
**Calculation of  $n$  and  $\text{deltax}$**   
**Loop over  $i$**

- Parallelize the code

### 1 Complete OpenMP statements

- Initialization
- omp get max threads
- omp get thread num

1 Parallelize the Loop over  $i$  :  
Use **omp parallel do/for**  
with the default(none) clause

2 Compile with:  
**icc -o pi -openmp c\_pi.c**  
Or  
**ifort -o pi -openmp f\_pi.F90**

3 Run with 1, 2, 4, 8, 12, 24 threads

e.g. export OMP\_NUM\_THREADS=4  
./c\_pi or ./f\_pi

4 **Compare timings**

- ✓ Timings decrease with more threads
- ✓ What is the scale up at 24 threads?.

# Exercise: neighbor update

- Double loop  $i,j$ : triangular iteration space
- Outer iterations do not take equal time
- Use schedule clause
  - Static schedule with small chunk size
  - Dynamic: has higher overhead
  - Guided: heuristic
- Control schedule through environment vars
- Also: reduction update

# Exercise: Neighbor Update; Part 1

- cd exercise\_neighbor
- Codes: f\_neighbor.f90/c\_neighbor.c

## neighbor update

### Parallel Region

Initialization: j\_update

### Parallelize loop i

Loop i

Loop j

increment j\_update

Loop k

b is calculated from a

- Try different schedules:  
static, dynamic, guided

Compile with: **make f\_neighbor**  
**make c\_neighbor**

- Parallelize the Loop over **i**
- Use a **single** construct for initialization
- Would a **master** construct work, too?
- Use **critical** for increment of **j\_update**
- Use omp parallel do/for with the default(none) clause

# Exercise: Neighbor Update; Part 2

## neighbor update

Parallel Region

Initialization: **j\_update**

Parallelize loop i

Loop i

Loop j

**single** or **master**

increment **j\_update**

**end single** or **end master**

Loop k

b is calculated from a

Compile with: **make f\_neighbor**  
**make c\_neighbor**

- Change the **single** to a **master** construct
- Run with 1 **and** 24 threads
- How does the number of **j\_update** change?

# Exercise: Red-Black Update; **Part 1**

- cd exercise\_redblack
- Codes: f\_red\_black.f90/c\_red\_black.c
- make a **copy** and create f\_red\_black\_v1.f90/c\_read\_black\_v1.c

## red-black update

Iteration Loop: **niter**

Loop: Update even elements

Loop: Update odd elements

Initialize **error**

Loop-summation: **error**

Compile with: **make f\_red\_black\_v1**  
**make c\_red\_black\_v1**

## **Part 1**

- Parallelize each loop separately
  - Use **omp parallel do/for** for the ‘ ‘Update’ ’ -loops
  - Use a **reduction** for the ‘ ‘Error’ ’ -calculation with the **default(none)** clause
- Try **static** scheduling

# Exercise: Red-Black Update; **Part 2**

- cd exercise\_redblack
- Start from **version 1**
- Codes: f\_red\_black.f90/c\_red\_black.c
- make a **copy** and create f\_red\_black\_v2.f90/c\_read\_black\_v2.c

## red-black update

Iteration Loop: **niter**

Loop:

Update even and odd el.

Initialize **error**

Loop-summation: **error**

- Try **static** scheduling

Compile with: **make f\_red\_black\_v2**  
**make c\_red\_black\_v2**

## **Part 1**

- Can the loops be combined?
- Why can the ‘update’ loops be combined?
- Why can the ‘error’ loop not be combined?
- Task:  
Combine the “update” loops

# Solution: Red-Black Update; **Part 2**

## red-black update

```
!*** Update even elements
do i=2, n, 2
    a(i) = 0.5 * (a(i) + a(i-1))
enddo
!*** Update odd elements
do i=1, n-1, 2
    a(i) = 0.5 * (a(i) + a(i+1))
enddo
```

## red-black update

```
!*** Update even and odd
!*** in one loop
do i=2, n, 2
    a(i) = 0.5 * (a(i) + a(i-1))
    a(i-1) = 0.5 * (a(i-1) + a(i))
enddo
```

# Exercise: Red-Black Update; **Part 3**

- cd exercise\_redblack
- Start from **version 2**
- Codes: f\_red\_black.f90/c\_red\_black.c
- make a **copy** and create f\_red\_black\_v3.f90/c\_read\_black\_v3.c

## red-black update

Iteration Loop: niter

**parallel region**

Loop:

Update even and odd el.

**single**

Initialize error

**end single**

Loop-summation: error

**end parallel region**

Compile with: **make f\_red\_black\_v3**  
**make c\_red\_black\_v3**

## **Part 1**

- Make **one parallel region** around both loops: “update” and “error”.
- The initialization of error has to be done by one thread
- Use a **single** construct
- Would a **master** construct work?



# Exercise: Orphaned work-sharing

- cd exercise\_print
- Codes: f\_print.f90/c\_print.c
- make a **copy** and create f\_red\_black\_v3.f90/c\_read\_black\_v3.c

## Orphaned work-sharing

parallel region

print 1

parallel Loop

print 2

call printer\_sub

master

print 5

subroutine print\_sub

parallel Loop

print 3

Loop

print 4

Compile with: **make f\_print**  
**make c\_print**

- Inspect the code
- Run with 1, 2, ... threads
- Explain the output
- How often are the 5 print statements executed?
- Why?

# Exercises – Lab 1

- Exercise 1: Kernel check  
f\_kernel.f90/c\_kernel.c  
Kernel of the calculation (see exercise 2)  
Parallelize one Loop
- Exercise 2: Calculation of  $\pi$   
f\_pi.f90/c\_pi.c  
Parallelize one Loop with a reduction
- Exercise 3: daxpy ( $a * x + b$ )  
f\_daxpy.f90/c\_daxpy.c  
Parallelize one Loop

# Exercises – Lab 2

- Exercise 4: Update from neighboring cells (2 arrays)  
f\_neighbor.f90/c\_neighbor.c  
Create a **Parallel Region**  
Use a **Single** construct to initialize  
Use a **Critical** construct to update  
Use **dynamic** or **guided** scheduling
- Exercise 5: Update from neighboring cells (same array)  
f\_red\_black.f90/c\_red\_black.c  
Parallelize 3 individual loops, use a reduction  
Create a **Parallel Region**  
Combine loops 1 and 2  
Use a **Single** construct to initialize