



OpenMP Thread Affinity

TACC OpenMP Team

milfeld/lars/agomez@tacc.utexas.edu

Learning Objective

-- Affinity --

- Understand why you should care about affinity
- Basic concepts
 - OMP_PLACES
 - OMP_PROC_BIND
- Standard affinity solutions for
 - Simple OpenMP
 - Nested OpenMP
 - Hybrid OpenMP (with MPI)
- Thread **and** Memory affinity

OpenMP Pre 3.0

- Affinity not controlled by OpenMP
- System default applied
- Other tools used

RaW = Read after Write, etc. 2

Why do we care about Thread Affinity?

Why do we want to control where threads are executed?

- Gain some (a little bit) of performance
- Prevent losing potentially a lot of performance

RaW = Read after Write, etc. 3

Why do we care about Thread Affinity?

Why do we want to control where threads are executed?

- Gain some (a little bit) of performance
- Prevent losing potentially a lot of performance

Good News

- It is very easy to do
- Basic setups will cover 95% of use cases

RaW = Read after Write, etc. 4

```

Tasks: 793 total,  2 running, 780 sleeping, 10 stopped,  1 zombie
Cpu0  : 69.3%us,  3.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu1  :  1.0%us,  6.9%sy,  0.0%ni, 92.1%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu2  : 73.3%us,  6.9%sy,  0.0%ni, 19.8%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu3  : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu4  : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu5  : 18.9%us,  1.1%sy,  0.0%ni, 80.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu6  : 68.6%us,  3.9%sy,  0.0%ni, 27.5%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu7  :  0.0%us,  1.0%sy,  0.0%ni, 99.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu8  : 71.0%us,  1.0%sy,  0.0%ni, 28.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu9  : 69.6%us,  2.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu10 : 70.3%us,  2.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu11 : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu12 : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu13 : 70.3%us,  2.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu14 : 70.3%us,  1.0%sy,  0.0%ni, 28.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu15 : 91.2%us,  1.0%sy,  0.0%ni,  6.9%id,  0.0%wa,  0.0%hi,  1.0%si,  0.0%st
Mem: 65922808k total, 60127364k used,  5795444k free, 35867560k buffers
Swap: 4194296k total,      0k used,  4194296k free, 15352588k cached

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
25125	lars	20	0	2628m	1.5g	1436	R	939.9	2.4	0:12.23	stencil.eo
24084	shempel	20	0	92852	3064	1236	S	24.1	0.0	0:11.44	sshd

Example from Stampede

- 16 cores → OMP_NUM_THREADS set to 16
- No pinning; uses default from the run-time
- Some cores are idling (Cpu1 and Cpu7)
- Some cores apparently oversubscribed

How much performance would you loose?

```

Tasks: 793 total,  2 running, 780 sleeping, 10 stopped,  1 zombie
Cpu0  : 69.3%us,  3.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu1  :  1.0%us,  6.9%sy,  0.0%ni, 92.1%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu2  : 73.3%us,  6.9%sy,  0.0%ni, 19.8%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu3  : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu4  : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu5  : 18.9%us,  1.1%sy,  0.0%ni, 80.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu6  : 68.6%us,  3.9%sy,  0.0%ni, 27.5%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu7  :  0.0%us,  1.0%sy,  0.0%ni, 99.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu8  : 71.0%us,  1.0%sy,  0.0%ni, 28.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu9  : 69.6%us,  2.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu10 : 70.3%us,  2.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu11 : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu12 : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu13 : 70.3%us,  2.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu14 : 70.3%us,  1.0%sy,  0.0%ni, 28.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu15 : 91.2%us,  1.0%sy,  0.0%ni,  6.9%id,  0.0%wa,  0.0%hi,  1.0%si,  0.0%st
Mem: 65922808k total, 60127364k used,  5795444k free, 35867560k buffers
Swap: 4194296k total,      0k used,  4194296k free, 15352588k cached

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
25125	lars	20	0	2628m	1.5g	1436	R	939.9	2.4	0:12.23	stencil.eo
24084	shempel	20	0	92852	3064	1236	S	24.1	0.0	0:11.44	sshd

Example from Stampede

- 16 cores → OMP_NUM_THREADS set to 16
- No pinning; uses default from the run-time
- Some cores are idling (Cpu1 and Cpu7)
- Some cores apparently oversubscribed

How much performance would you loose?

- **Dynamic scheduling: 2 out of 16 = 12.5%**

```

Tasks: 793 total,  2 running, 780 sleeping,  10 stopped,  1 zombie
Cpu0  : 69.3%us,  3.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu1  :  1.0%us,  6.9%sy,  0.0%ni, 92.1%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu2  : 73.3%us,  6.9%sy,  0.0%ni, 19.8%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu3  : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu4  : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu5  : 18.9%us,  1.1%sy,  0.0%ni, 80.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu6  : 68.6%us,  3.9%sy,  0.0%ni, 27.5%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu7  :  0.0%us,  1.0%sy,  0.0%ni, 99.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu8  : 71.0%us,  1.0%sy,  0.0%ni, 28.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu9  : 69.6%us,  2.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu10 : 70.3%us,  2.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu11 : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu12 : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu13 : 70.3%us,  2.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu14 : 70.3%us,  1.0%sy,  0.0%ni, 28.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu15 : 91.2%us,  1.0%sy,  0.0%ni,  6.9%id,  0.0%wa,  0.0%hi,  1.0%si,  0.0%st
Mem: 65922808k total, 60127364k used,  5795444k free, 35867560k buffers
Swap: 4194296k total,      0k used,  4194296k free, 15352588k cached

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
25125	lars	20	0	2628m	1.5g	1436	R	939.9	2.4	0:12.23	stencil.eo
24084	shempel	20	0	92852	3064	1236	S	24.1	0.0	0:11.44	sshd

Example from Stampede

- 16 cores → OMP_NUM_THREADS set to 16
- No pinning; uses default from the run-time
- Some cores are idling (Cpu1 and Cpu7)
- Some cores apparently oversubscribed

How much performance would you loose?

- Dynamic scheduling: 2 out of 16 = 12.5%
- **Static scheduling: 50% (waiting for the oversubscribed cores to finish)**


```

Tasks: 793 total,  2 running, 780 sleeping, 10 stopped,  1 zombie
Cpu0  : 69.3%us,  3.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu1  :  1.0%us,  6.9%sy,  0.0%ni, 92.1%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu2  : 73.3%us,  6.9%sy,  0.0%ni, 19.8%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu3  : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu4  : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu5  : 18.9%us,  1.1%sy,  0.0%ni, 80.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu6  : 68.6%us,  3.9%sy,  0.0%ni, 27.5%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu7  :  0.0%us,  1.0%sy,  0.0%ni, 99.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu8  : 71.0%us,  1.0%sy,  0.0%ni, 28.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu9  : 69.6%us,  2.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu10 : 70.3%us,  2.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu11 : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu12 : 70.6%us,  1.0%sy,  0.0%ni, 28.4%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu13 : 70.3%us,  2.0%sy,  0.0%ni, 27.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu14 : 70.3%us,  1.0%sy,  0.0%ni, 28.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu15 : 91.2%us,  1.0%sy,  0.0%ni,  6.9%id,  0.0%wa,  0.0%hi,  1.0%si,  0.0%st
Mem: 65922808k total, 60127364k used,  5795444k free, 35867560k buffers
Swap: 4194296k total,      0k used,  4194296k free, 15352588k cached

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
25125	lars	20	0	2628m	1.5g	1436	R	939.9	2.4	0:12.23	stencil.eo
24084	shempel	20	0	92852	3064	1236	S	24.1	0.0	0:11.44	sshd

Example from Stampede

- 16 cores → OMP_NUM_THREADS set to 16
- No pinning; uses default from the run-time
- Some cores are idling (Cpu1 and Cpu7)
- Some cores apparently oversubscribed

How much performance would you loose?

- Dynamic scheduling: 2 out of 16 = 12.5%
- Static scheduling: 50% (waiting for the oversubscribed cores to finish)
- Implications for MPI codes even more severe

How to use 'top' to look at core utilization

- Open a terminal
- Type **top**
- Press the 1 key (#1 key top left corner)
- **top** may complain about the window size
- You may have to increase the window so that there are more lines in the window than cores
- You may increase the update frequency. Type 's' and 1, which will change the update interval to 1 second

We will use 2 windows in the lab

One to run the experiment

The other to inspect core utilization

Terminology

- Compute node: basic building block of a cluster
 - A single desktop is also a node
- Socket: each node may have 1, 2, or 4 sockets
 - Sockets are often called CPUs (confusion!)
- Cores: each socket may have any (small) number of cores; 2, 4, 8, 12, 16, 20
 - Cores are sometimes called CPUs (more confusion)
- Each core may have 1 or 2 (Hyperthreading) hardware threads (hw threads)
 - Some architectures have 4 hw threads
 - Lonestar has Hyperthreading enabled
- OpenMP threads use the hardware threads

10

Typical Thread Utilization

All cores are being used

- Without Hyperthreading

- 1 OpenMP thread per hardware thread
- 1 OpenMP thread per core

- With Hyperthreading

- 1 OpenMP thread per hardware thread
- 2 OpenMP threads per core

Or

- 1 OpenMP thread per core
- Every other hardware thread is not used = Hyperthreading enabled but not utilized

Typical Thread Utilization

Exceptions

- Oversubscribing one core
 - 16 cores and 17 threads
 - Extra thread doing different work
 - Scheduling, MPI communication, or I/O, etc.
- Leaving cores idle
 - Memory bandwidth limited code
 - Increasing shared L3 cache per thread
- Experiment to find best performance
 - One thread per core
 - One thread per hardware thread
 - ‘Odd’ configurations

12

How to pin OpenMP threads to a resource

OMP_PLACES and **OMP_PROC_BIND**

- **Step 1: Defining OpenMP Places on a node**
 - Environment variable: **OMP_PLACES**
 - Default: 1 node is one big place
 - Divide node in multiple places
 - Example: 2 sockets, each a place (2 places total)
 - or: each core a separate place (16 places on Stampede)
 - Environment variable: **OMP_PLACES**
- **Step 2: Pin OpenMP threads (procs) to the Places**
 - Environment variable: **OMP_PROC_BIND**
 - Pinning can be 'spread out' among places
 - Pinning can be 'close together'
- **Places and pinning are controlled by environment variables**

Defining OpenMP Places: **OMP_PLACES**

Example without Hyperthreading

- Each core is 1 hardware thread
- The cores are numbered
- Stampede: 16 cores, numbered from 0 to 15
- Example for 2 places on a node
 - `export OMP_PLACES="{0,1,2,3,4,5,6,7}{8,9,10,11,12,13,14,15}"`
 - First place: Cores 0-7 (socket 0)
 - Second place: Cores 8-15 (socket 1)
- Example for 16 places
 - `Export OMP_PLACES="{0}{1}{2}{3}{4}{5}{6}{7}{8}{9}{10}{11}{12}{13}{14}{15}"`
- Syntax shortcuts available (not shown, buggy at times)

Pinning Policy: `OMP_PROC_BIND`

Example without Hyperthreading

- Chose between '**spread**' and '**close**'
 - **Spread**: Choose places far away from each other
 - **Close**: Choose places close together
 - A third possibility is 'master'

Let's look at some examples

- Different combinations of
 - `OMP_NUM_THREADS`
 - `OMP_PLACES`
 - `OMP_PROC_BIND`

Example 1 (No Hyperthreading)

Stampede: 16 cores, 16 Places

- export OMP_NUM_THREADS=4
- export OMP_PLACES="{0}{1}{2}{3}{4}{5}{6}{7}{8}{9}{10}{11}{12}{13}{14}{15}"
- export OMP_PROC_BIND=spread
- Threads may be bound to cores 0, 4, 8, and 12 or 1, 5, 9, and 13, etc.
 - ‘spread’ = maximum distance is 4 (16 places, 4 threads)

Example 1

```
Tasks: 825 total, 2 running, 812 sleeping, 10 stopped, 1 zombie
Cpu0  :100.0%us, 0.0%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu1  : 0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu2  : 0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu3  : 1.9%us, 0.0%sy, 0.0%ni, 98.1%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu4  : 96.2%us, 3.8%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu5  : 1.0%us, 2.9%sy, 0.0%ni, 96.2%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu6  : 0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu7  : 0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu8  : 97.1%us, 2.9%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu9  : 0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu10 : 0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu11 : 0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu12 : 96.2%us, 3.8%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu13 : 0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu14 : 0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu15 : 0.0%us, 2.9%sy, 0.0%ni, 88.2%id, 7.8%wa, 0.0%hi, 1.0%si, 0.0%st
Mem: 65922808k total, 52179032k used, 13743776k free, 35169356k buffers
Swap: 4194296k total, 0k used, 4194296k free, 9017168k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
7860	lars	20	0	1812m	1.5g	1448	R	400.6	2.4	0:08.11	stencil.eo
6172	cazes	20	0	152m	10m	1420	S	4.8	0.0	51:40.45	wget

17

Example 2 (No Hyperthreading)

Stampede: 16 cores, 16 Places

- export OMP_NUM_THREADS=4
- export OMP_PLACES="{0}{1}{2}{3}{4}{5}{6}{7}{8}{9}{10}{11}{12}{13}{14}{15}"
- export OMP_PROC_BIND=close
- Threads may be bound to cores 0, 1, 2, and 3 or 1, 2, 3, 4, etc.
 - 'close' = minimum distance is 1

Example 2

```
Tasks: 825 total, 2 running, 812 sleeping, 10 stopped, 1 zombie
Cpu0  : 96.0%us, 2.8%sy, 0.0%ni, 1.2%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu1  : 72.8%us, 2.5%sy, 0.0%ni, 24.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu2  : 73.5%us, 1.5%sy, 0.0%ni, 25.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu3  : 75.7%us, 6.5%sy, 0.0%ni, 17.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu4  : 3.7%us, 2.8%sy, 0.0%ni, 93.6%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu5  : 0.6%us, 0.9%sy, 0.0%ni, 98.5%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu6  : 0.0%us, 0.3%sy, 0.0%ni, 99.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu7  : 0.6%us, 0.9%sy, 0.0%ni, 98.5%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu8  : 4.0%us, 2.2%sy, 0.0%ni, 93.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu9  : 0.3%us, 0.3%sy, 0.0%ni, 99.4%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu10 : 0.0%us, 0.3%sy, 0.0%ni, 99.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu11 : 0.0%us, 0.3%sy, 0.0%ni, 99.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu12 : 0.0%us, 0.0%sy, 0.0%ni, 100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu13 : 0.3%us, 0.0%sy, 0.0%ni, 99.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu14 : 0.0%us, 0.0%sy, 0.0%ni, 100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu15 : 0.3%us, 3.7%sy, 0.0%ni, 93.5%id, 2.5%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 65922808k total, 53083312k used, 12839496k free, 35169388k buffers
Swap: 4194296k total, 0k used, 4194296k free, 9739544k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
7914	lars	20	0	1812m	1.5g	1456	R	342.6	2.4	0:13.17	stencil.eo
6172	cazes	20	0	152m	10m	1420	S	4.7	0.0	51:42.11	wget

19

Example 3 (No Hyperthreading)

Stampede: 16 cores, 2 Places

- export OMP_NUM_THREADS=4
- export OMP_PLACES="{0,1,2,3,4,5,6,7}{8,9,10,11,12,13,14,15}"
- export OMP_PROC_BIND=s~~pread~~
- 2 threads will be bound to the first place (threads 0 and 2)
- 2 threads will be bound to the second place (threads 1 and 3)
- Within a place threads can move around, so
 - Threads 0 and 2 may be executed by any core 0-7
 - Threads 1 and 3 may be executed on any core 8-15

Example 3

```
Tasks: 823 total, 2 running, 810 sleeping, 10 stopped, 1 zombie
Cpu0  :  0.0%us,  0.0%sy,  0.0%ni,100.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu1  : 99.0%us,  1.0%sy,  0.0%ni,  0.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu2  :  0.0%us,  0.0%sy,  0.0%ni,100.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu3  :  1.0%us,  1.9%sy,  0.0%ni, 97.1%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu4  :  0.0%us,  0.0%sy,  0.0%ni,100.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu5  : 83.3%us,  2.9%sy,  0.0%ni, 13.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu6  :  0.0%us,  0.0%sy,  0.0%ni,100.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu7  :  0.0%us,  0.0%sy,  0.0%ni,100.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu8  : 83.3%us,  2.9%sy,  0.0%ni, 13.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu9  :  0.0%us,  0.0%sy,  0.0%ni,100.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu10 : 83.3%us,  2.9%sy,  0.0%ni, 13.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu11 :  0.0%us,  0.0%sy,  0.0%ni,100.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu12 :  0.0%us,  0.0%sy,  0.0%ni,100.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu13 :  0.0%us,  0.0%sy,  0.0%ni,100.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu14 :  0.0%us,  0.0%sy,  0.0%ni,100.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Cpu15 :  1.0%us,  1.0%sy,  0.0%ni, 98.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Mem: 65922808k total, 52366168k used, 13556640k free, 35169460k buffers
Swap: 4194296k total, 0k used, 4194296k free, 9164988k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
7991	lars	20	0	1812m	1.5g	1452	R	359.5	2.4	0:14.38	stencil.eo
6172	cazes	20	0	152m	10m	1420	S	2.0	0.0	51:46.22	wget

21

Example 4 (No Hyperthreading)

Stampede: 16 cores, 16 Places

- export OMP_NUM_THREADS=16
- export OMP_PLACES="{0}{1}{2}{3}{4}{5}{6}{7}{8}{9}{10}{11}{12}{13}{14}{15}"
- export OMP_PROC_BIND=spread
- Threads may be bound to cores 0, 4, 8, and 12 or 1, 5, 9, and 13, etc.
 - ‘spread’ = maximum distance is 1 (16 places, 16 threads)
 - So each thread is bound to a different place
- This is a standard situation: covering 95% of the applications
 - 1 thread per core, no ‘overlap’
 - Threads cannot move around

Example 4

```
Tasks: 826 total, 3 running, 812 sleeping, 10 stopped, 1 zombie
Cpu0  : 99.0%us, 1.0%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu1  : 89.3%us, 10.7%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu2  : 92.1%us, 3.0%sy, 0.0%ni, 5.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu3  : 88.1%us, 6.9%sy, 0.0%ni, 5.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu4  : 88.3%us, 5.8%sy, 0.0%ni, 5.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu5  : 89.2%us, 4.9%sy, 0.0%ni, 5.9%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu6  : 92.2%us, 1.9%sy, 0.0%ni, 5.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu7  : 88.3%us, 5.8%sy, 0.0%ni, 5.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu8  : 90.3%us, 3.9%sy, 0.0%ni, 5.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu9  : 85.4%us, 8.7%sy, 0.0%ni, 5.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu10 : 91.3%us, 2.9%sy, 0.0%ni, 5.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu11 : 89.3%us, 4.9%sy, 0.0%ni, 5.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu12 : 90.3%us, 3.9%sy, 0.0%ni, 5.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu13 : 78.2%us, 16.8%sy, 0.0%ni, 5.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu14 : 90.3%us, 3.9%sy, 0.0%ni, 5.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu15 : 87.3%us, 6.9%sy, 0.0%ni, 5.9%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 65922808k total, 53124472k used, 12798336k free, 35169488k buffers
Swap: 4194296k total, 0k used, 4194296k free, 9772000k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
8029	lars	20	0	2628m	1.5g	1452	R	1453.4	2.4	0:57.95	stencil.eo
8024	root	20	0	108m	952	408	R	53.0	0.0	0:02.73	rsync

23

Standard Setup (95% of cases)

Stampede: 16 cores, 16 places, 16 threads

- export OMP_NUM_THREADS=16
- export OMP_PLACES=cores
- export OMP_PROC_BIND=spread
- Every OpenMP thread is pinned to a core/hardware thread

Lonestar: 24 cores, 24 places, 24 or 48 threads

- export OMP_NUM_THREADS=24
- export OMP_NUM_THREADS=48
- export OMP_PLACES=cores (or threads)
- export OMP_PROC_BIND=spread
- One or two OpenMP thread(s) pinned to a core with 2 hardware threads

24

Summary OpenMP code

No Hyperthreading

- export OMP_NUM_THREADS=<number of cores>
- export OMP_PLACES=cores (shortcut for a list of cores)
- export OMP_PROC_BIND=spread

With Hyperthreading

- export OMP_NUM_THREADS=<1 or 2 x number of cores>
 - export OMP_PLACES=cores
 - export OMP_PROC_BIND=spread
 - Each thread has its own core
 - With Hyperthreading 2 threads are sharing a core
-
- Increase/reduce number of OpenMP threads if you want to oversubscribe cores or want to leave cores idle

Nested OpenMP (1)

Example: 2-level nested OpenMP

- Upper level with 2 threads (Outer parallel region)
- Lower level with 4 threads (Inner parallel region)

You spawn 2 OpenMP threads

Each thread spawns 4 sub-threads

Good configuration would be to start the 2 upper-level threads as far away as possible. This would be one thread per socket.

And then use the rest of the socket to spawn the lower-level threads

Nested OpenMP (2)

```
export OMP_PLACES=cores  
export OMP_PROC_BIND=spread,close  
export OMP_NUM_THREADS=4
```

OMP_NUM_THREADS applies to both levels. To change the number of threads in the upper level to 2 I hard-wired this in the code. See example in the lab later.

OMP_PROC_BIND accepts multiple arguments

First argument: ^Spread is for the outer level

Second argument: ^Close is for the lower level

Example Nested (close)

```
Tasks: 592 total, 2 running, 590 sleeping, 0 stopped, 0 zombie
Cpu0  :100.0%us, 0.0%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu1  : 88.2%us, 2.0%sy, 0.0%ni, 9.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu2  : 87.1%us, 3.0%sy, 0.0%ni, 9.9%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu3  : 89.1%us, 1.0%sy, 0.0%ni, 9.9%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu4  :  0.0%us, 0.0%sy, 0.0%ni, 98.0%id, 2.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu5  :  0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu6  :  0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu7  :  0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu8  : 70.3%us, 3.0%sy, 0.0%ni, 26.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu9  : 70.6%us, 2.0%sy, 0.0%ni, 27.5%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu10 : 71.3%us, 2.0%sy, 0.0%ni, 26.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu11 : 72.3%us, 1.0%sy, 0.0%ni, 26.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu12 :  1.0%us, 1.0%sy, 0.0%ni, 98.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu13 :  0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu14 :  0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu15 :  0.0%us, 0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 32815324k total, 5520460k used, 27294864k free, 3484k buffers
Swap:  0k total,  0k used,  0k free, 35388k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
124736	lars	20	0	2084m	1.5g	1460	R	661.6	4.8	0:25.34	stencil.eo
243	root	39	19	0	0	0	S	1.0	0.0	104:03.51	kipmi0

28

Hybrid: OpenMP and MPI

```
export OMP_PLACES=cores  
export OMP_PROC_BIND=sred  
export OMP_NUM_THREADS=4
```

Job launched with ibrun and tacc_affinity

Example is with 2 MPI tasks on a node; one per socket

```
lbrun tacc_affinity a.out
```

Example Nested (spread)

```
Tasks: 599 total, 3 running, 596 sleeping, 0 stopped, 0 zombie
Cpu0  : 67.3%us, 32.7%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu1  :  0.0%us,  0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu2  : 32.4%us,  1.0%sy, 0.0%ni, 66.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu3  :  0.0%us,  0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu4  : 32.4%us,  1.0%sy, 0.0%ni, 66.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu5  :  0.0%us,  0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu6  : 32.4%us,  1.0%sy, 0.0%ni, 66.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu7  :  0.0%us,  0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu8  : 65.0%us, 35.0%sy, 0.0%ni,  0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu9  :  0.0%us,  1.0%sy, 0.0%ni, 99.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu10 : 28.4%us,  1.0%sy, 0.0%ni, 70.6%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu11 :  0.0%us,  0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu12 : 29.7%us,  0.0%sy, 0.0%ni, 70.3%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu13 :  0.0%us,  0.0%sy, 0.0%ni,100.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu14 : 29.4%us,  0.0%sy, 0.0%ni, 70.6%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu15 :  0.0%us,  1.0%sy, 0.0%ni, 99.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 32815324k total, 7196848k used, 25618476k free, 5248k buffers
Swap:  0k total,  0k used,  0k free, 121832k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
126445	lars	20	0	1628m	1.5g	5444	R	199.5	4.8	0:03.65	stencil.eo
126446	lars	20	0	1627m	1.5g	5408	R	187.7	4.8	0:03.53	stencil.eo

30

Hybrid: OpenMP and MPI

```
export OMP_PLACES=cores  
export OMP_PROC_BIND=close  
export OMP_NUM_THREADS=4
```

Job launched with ibrun and tacc_affinity

Example is with 2 MPI tasks on a node; one per socket

```
lbrun tacc_affinity a.out
```


Example Hybrid

```
Tasks: 602 total,   3 running, 599 sleeping,   0 stopped,   0 zombie
Cpu0  : 99.0%us,   1.0%sy,   0.0%ni,   0.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu1  : 98.0%us,   2.0%sy,   0.0%ni,   0.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu2  : 98.0%us,   2.0%sy,   0.0%ni,   0.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu3  : 98.0%us,   2.0%sy,   0.0%ni,   0.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu4  :  0.0%us,   2.0%sy,   0.0%ni,  95.1%id,   2.9%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu5  :  0.0%us,   1.0%sy,   0.0%ni,  99.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu6  :  0.0%us,   0.0%sy,   0.0%ni, 100.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu7  :  0.0%us,   0.0%sy,   0.0%ni, 100.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu8  :100.0%us,   0.0%sy,   0.0%ni,   0.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu9  : 97.1%us,   2.9%sy,   0.0%ni,   0.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu10 : 96.1%us,   3.9%sy,   0.0%ni,   0.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu11 : 99.0%us,   1.0%sy,   0.0%ni,   0.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu12 :  0.0%us,   0.0%sy,   0.0%ni, 100.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu13 :  0.0%us,   0.0%sy,   0.0%ni, 100.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu14 :  0.0%us,   0.0%sy,   0.0%ni, 100.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Cpu15 :  0.0%us,   0.0%sy,   0.0%ni, 100.0%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Mem: 32815324k total, 7194504k used, 25620820k free,    5080k buffers
Swap:      0k total,      0k used,      0k free, 121604k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
125925	lars	20	0	1628m	1.5g	5444	R	400.9	4.8	0:15.32	stencil.eo
125926	lars	20	0	1627m	1.5g	5404	R	399.9	4.8	0:14.98	stencil.eo

32

Pinning within the code

- More flexibility
- Different pinning strategies for different parallel regions
- Clause **proc_bind(*arg*)** added to parallel region
- Note: OpenMP places can only be defined once before the execution start

C/C++

```
#pragma omp parallel proc_bind(spread)
```

Fortran

```
!$omp parallel proc_bind(spread)
```

Summary

- Use 3 OpenMP environment variables
 - OMP_PLACES
 - OMP_PROC_BIND
 - OMP_NUM_THREADS
 - Binding and number of threads may also be change in code
- Choose places and binding so that the threads are on dedicated