

Introduction to Programming with OpenMP

OpenMP Team

Kent Milfeld, Lars Koesterke, and Antonio Gomez

milfeld/lars/agomez@tacc.utexas.edu

2016/03/31

Sides at: tinyurl.com/tacc-openmp or portal.tacc.utexas.edu/training click View Details

Outline

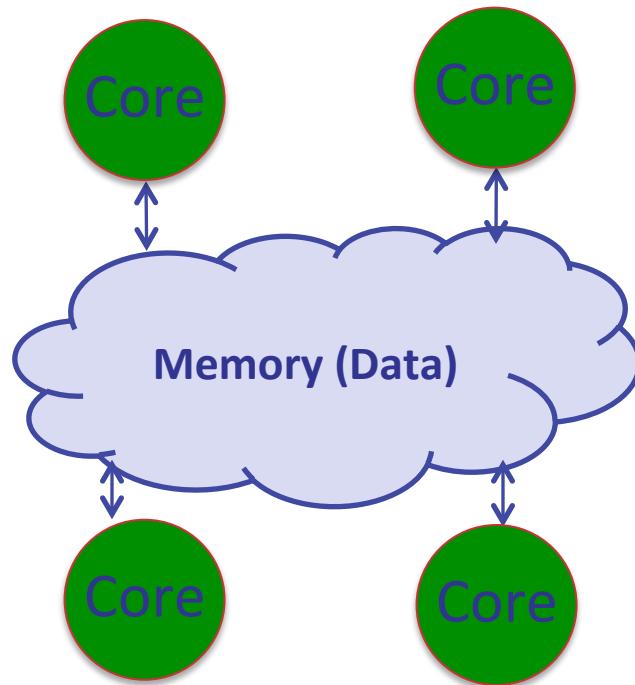
- What is OpenMP?
- How OpenMP Parallelism Works
 - Architecture
 - Parallelism
 - Execution Model (Fork-Join)
 - Memory Model & Synchronization
- OpenMP Syntax and API
 - Compiler Directives
 - Runtime Library Routines
 - Environment variables
- OpenMP 3.0 Changes

What is OpenMP?

- OpenMP stands for **Open Multi-Processing**
- A directive-based programming language for executing program components in parallel within a shared memory architecture.
- Three primary components are:
 - Directives to the Compiler within code (pragmas in C/C++, comments in Fortran)
 - Runtime Library Routines
 - Environment Variables
- Standard specifies C, C++, and FORTRAN Directives & API
- <http://www.openmp.org/> has the specification, examples, tutorials and documentation

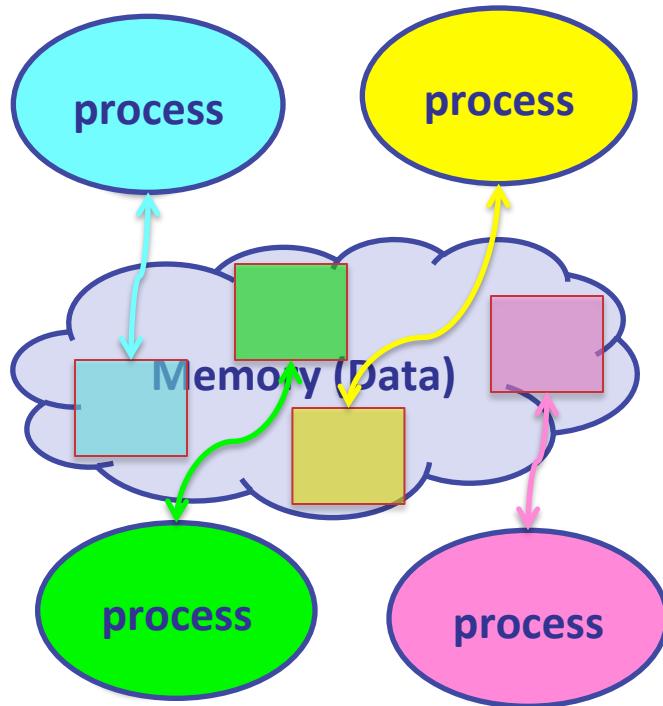
Hardware Architecture

- Cores: where processes or lightweight processes (threads) can run.
- Memory
 - Multiple cores have access to the Memory
 - Processes have their own section of memory to use



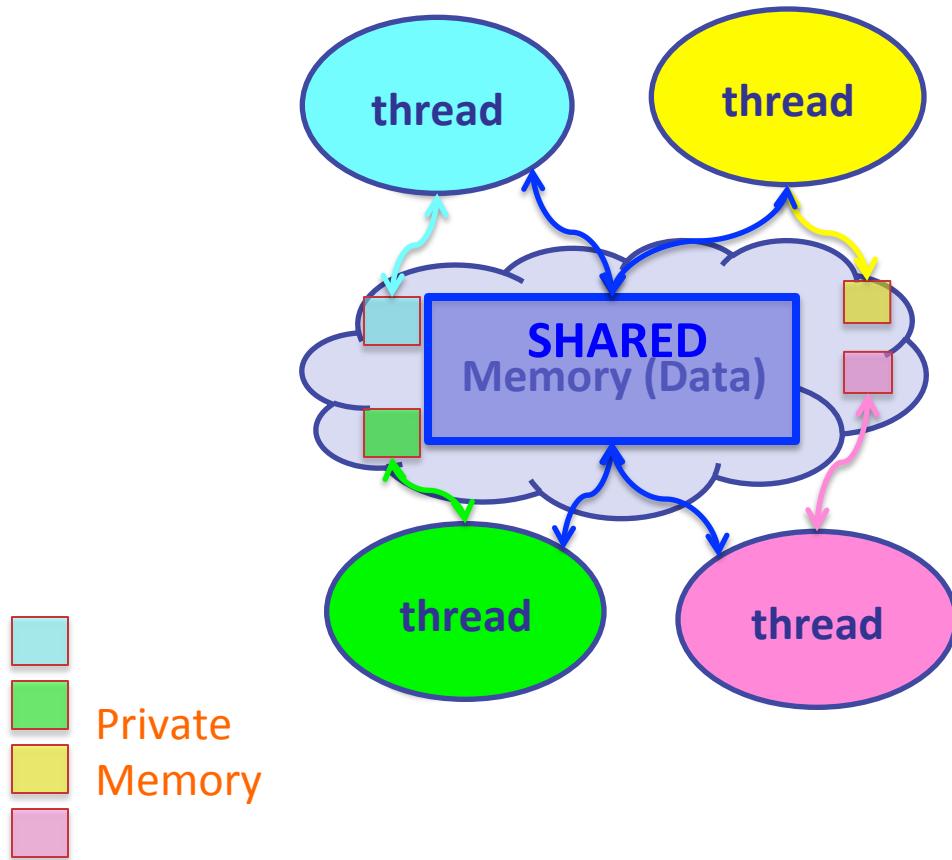
Software Architecture

- Cores: where processes or lightweight processes (threads) can run.
- Memory
 - Multiple cores have access to the Memory
 - Processes have their own section of memory to use



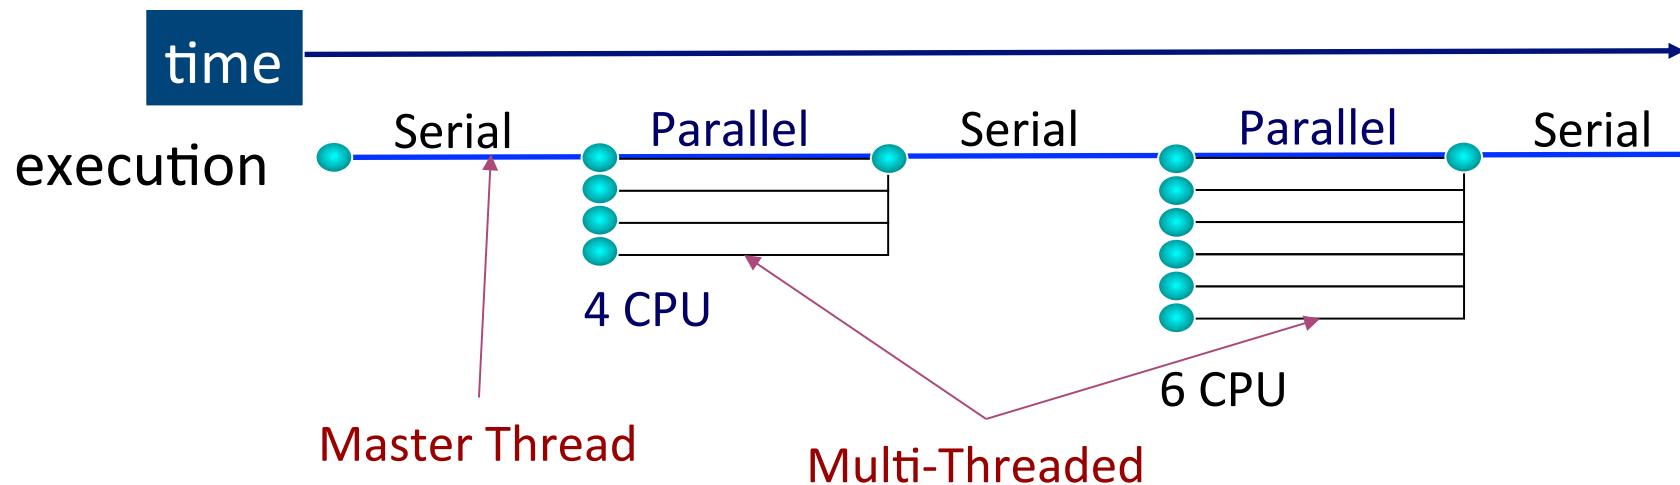
Software Architecture

- Threads Execute on the Cores
- Threads within an application can share memory.
- Data: shared or private
 - Shared data: all threads can access data in shared memory
 - Private data: can only be accessed by threads that own it



OpenMP Fork-Join Parallelism

- Programs begin as a single process: master thread
- Master thread executes in serial mode until the parallel region construct is encountered
- After executing the statements in the parallel region, team threads synchronize and terminate (join) but master continues



Memory Model and Synchronization

- Every thread has access to shared memory.
When multiple threads are created all threads share the same address space.
- Simultaneous updates to shared memory can create a *race condition*. Results change with different thread scheduling.
- Use mutual exclusion directives to avoid race conditions - but don't use too many because this will serialize performance.
- Barriers are used to synchronize threads.
- A flush exists at barriers, to make local changes to data (writes) seen by all threads.

Race conditions

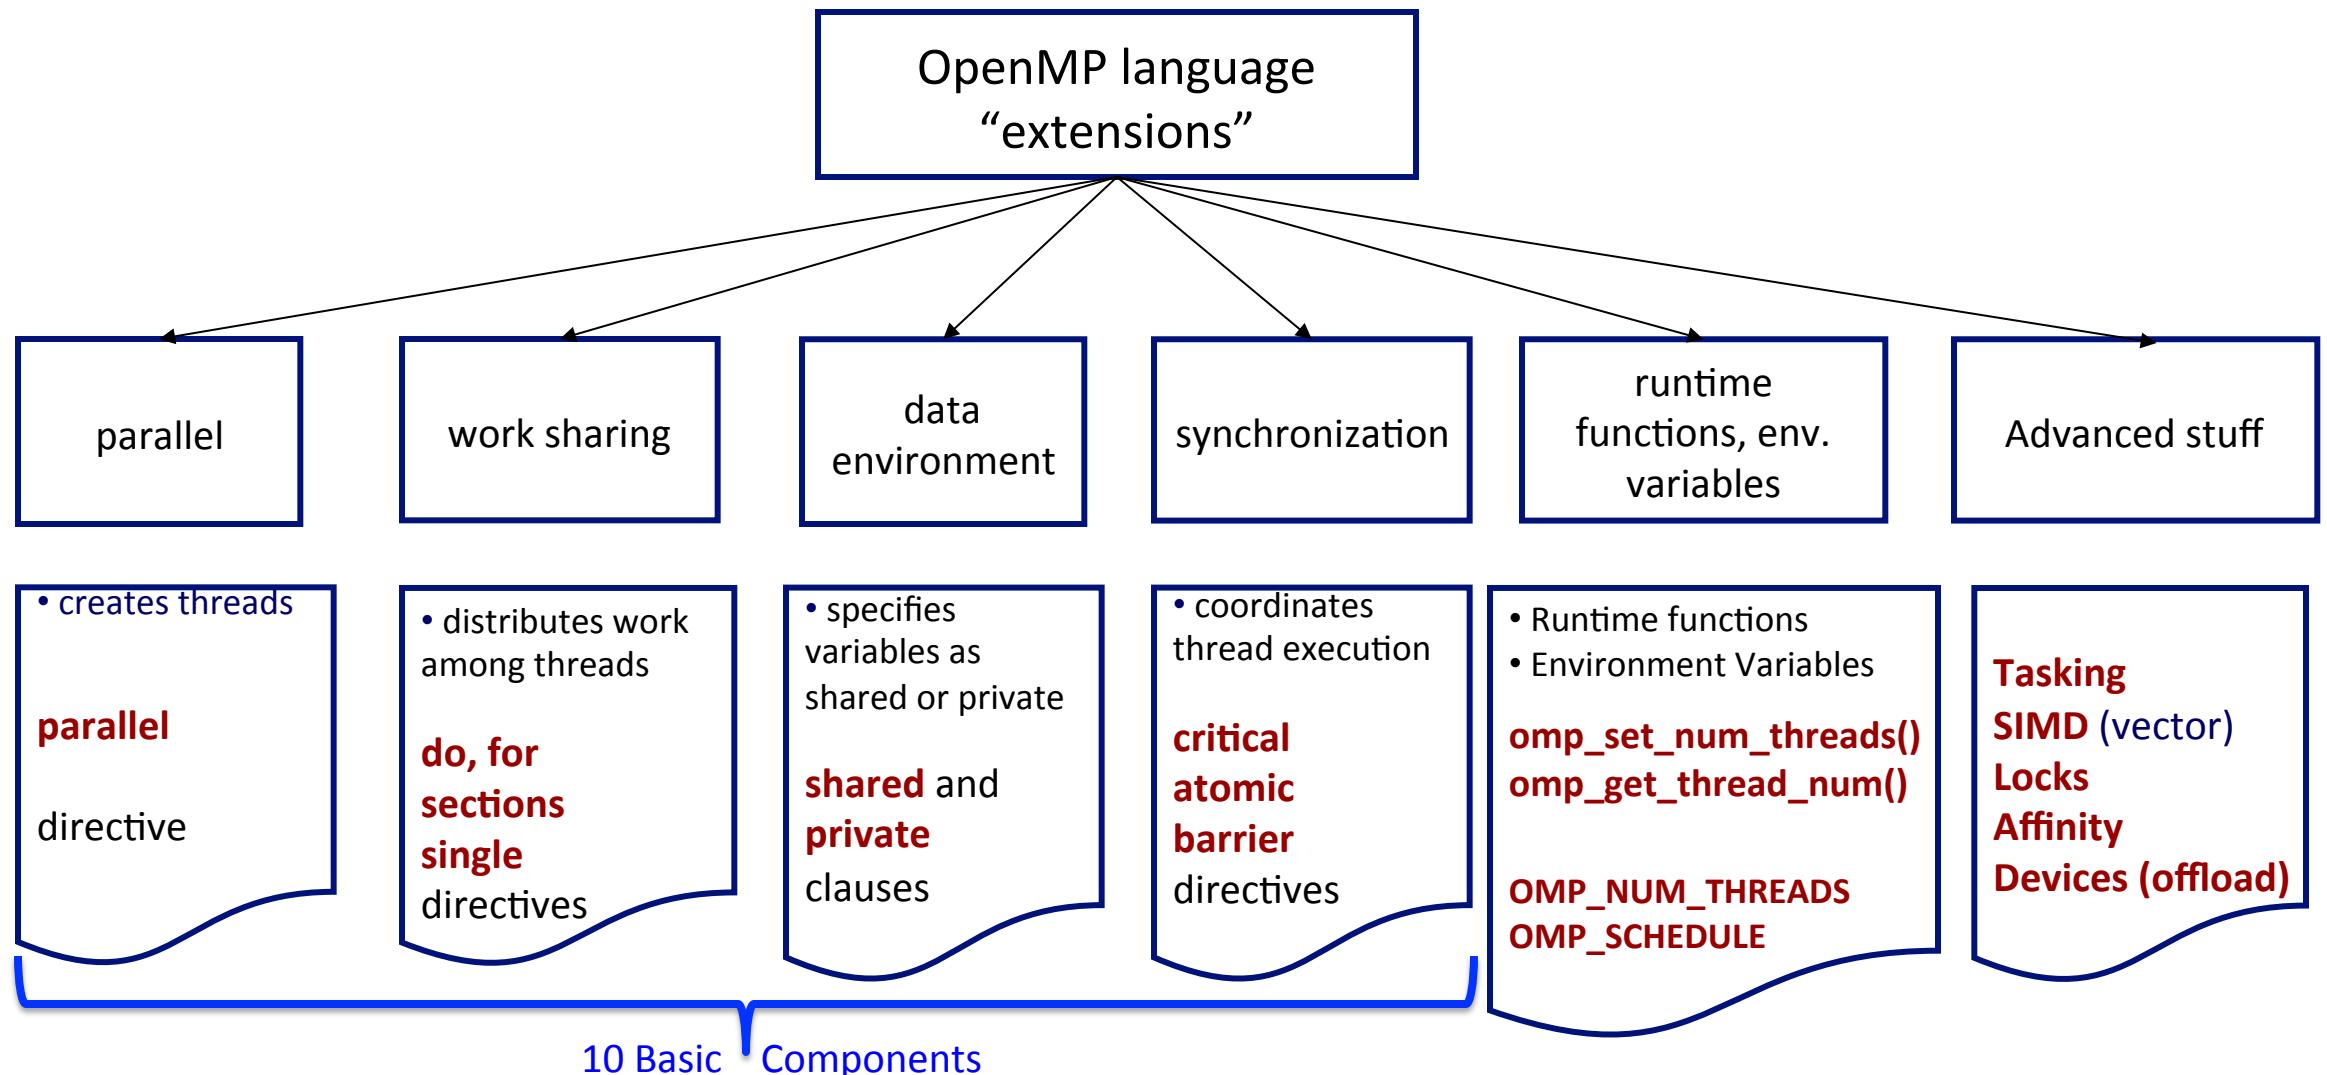
Thread 0: $i = i+2$

Thread 1: $i = i+3$

Initial value of i in memory = 0.

scenario 1.		scenario 2.		scenario 3.	
Thread 0	Thread 1	Thread 0	I = 0	Thread 1	Thread 0
read I = 0					
compute I = 2	compute I = 3	compute I = 2	compute I = 3	compute I = 3	compute I = 2
write I = 2	write I = 3		write I = 2		write I = 2
					read I = 2
					compute I = 5
					write I = 5
I = 3		I = 2		I = 5	

OpenMP Constructs



OpenMP Syntax

- Function prototypes and types are in the files:

`#include <omp.h>`

C

`use omp_lib`

F90

- Compiler directive syntax:

`#pragma omp construct [clause [[,]clause]...]`

C

`!$omp construct [clause [[,]clause]...]`

F90

- Example

`#pragma omp parallel num_threads(4)`

C

`!$omp parallel num_threads(4)`

F90

- Most OpenMP constructs apply to a “structured block”; that is, a block of one or more statements with one point of entry at the top and one point of exit at the bottom

OpenMP Directive Scope

- OpenMP directives are comments/pragmas in source:
 - F90 : **!\$OMP** free-format*
 - C/C++ : directives begin with the **# pragma omp** sentinel

Parallel regions are marked by enclosing parallel directives

F90 : **parallel ... end parallel**

C/C++ : Enclosed in block **{ ... }**, or single statement

- Work-sharing loops are marked by parallel do/for

* Fortran Fixed Format:
!\$OMP, **C\$OMP** or ***\$OMP**

Fortran

```
!$omp parallel  
...  
!$omp end parallel  
  
!$omp parallel  
  call foo(...)  
!$omp end parallel
```

C/C++

```
# pragma omp parallel  
{...}  
  
# pragma omp parallel for  
  foo(...);
```

Parallel Region & Work-Sharing

Use OpenMP directives to specify Parallel Region & Work-Sharing constructs

```
parallel  
[  
    do / for  
    sections  
    single  
]  
end parallel
```

Code block

do / for
sections
single

Each Thread Executes

Work Sharing
Work Sharing
Work Sharing (one thread)

Sentinels
“!\$omp” and
“#pragma omp”) not shown here.

Parallel Regions

```
1 #pragma omp parallel  
2 {  
3     code block  
4     work(...);  
5 }
```

- Line 1 Team of threads formed at parallel region
- Lines 3-4 Each thread executes code block and subroutine calls.
No branching (in or out) in a parallel region
- Line 5 All threads synchronize at end of parallel region (implied barrier)

Use the thread number to divide work among threads

Parallel Regions

```
1 !$OMP PARALLEL  
2   code block  
3   call work(...)  
4 !$OMP END PARALLEL
```

Line 1 Team of threads formed at parallel region.

Lines 2-3 Each thread executes code block and subroutine calls. No branching (in or out) in a parallel region.

Line 4 All threads synchronize at end of parallel region (implied barrier).

Use the thread number to divide work among.

Parallel Region & Number of Threads

- For example, to create a 4-thread Parallel region.
- export OMP_NUM_THREADS=4; a.out

```
double A[1000];  
  
#pragma omp parallel  
{  
    int ID = omp_get_thread_num();  
    foo(ID, A);  
}
```

In C++ a local variables are private to the thread...

- Each thread executes the code within the structured block
- Thread numbers range from 0 to Nthreads-1
- Each thread calls foo(ID,A) with a different ID {= 0 to 3}

Parallel Region & Number of Threads

- For example, to create a 4-thread Parallel region.
- export OMP_NUM_THREADS=4; a.out

```
real :: A(1000; integer :: ID  
!  
!$omp parallel  
  
    ID = omp_get_thread_num()  
    call foo(ID, A)  
  
!$omp end parallel
```

But we need to make id
Private to the thread– more later

- Each thread executes the code within the structured block
- Thread numbers range from 0 to Nthreads-1
- Each thread calls foo(ID,A) with a different ID {= 0 to 3}

Work sharing: Loop

```
1 #pragma omp parallel
2 {
3     #pragma omp for
4     for (i=0; i<N; i++)
5     {
6         a[i] = b[i] + c[i];
7     }
8 }
```

Or

```
#pragma omp parallel
#pragma omp for
for (i=0; i<N; i++)
    a[i] = b[i] +
c[i];
```

Line 1 Team of threads formed (parallel region).

Line 3-7 Loop iterations are split among threads.
implied barrier at }

Each loop iteration must be independent of other iterations.

Work sharing: Loop

```
1 !$OMP PARALLEL
2   !$OMP DO
3     do i=1,N
4       a(i) = b(i) + c(i)
5     enddo
6   !$OMP END PARALLEL
```

Line 1 Team of threads formed (parallel region).

Line 3-4 Loop iterations are split among threads by DO construct.

Line 5 !\$OMP END DO is optional after the enddo.

Line 5 Implied barrier at enddo.

Each loop iteration must be independent of other iterations.

OpenMP Combined Constructs

Fortran

```
!$omp parallel
  !$omp do
    do ...
    end do
  !$omp end do
 !$omp end parallel
```

```
!$omp parallel do
  do ...
  end do
```

!*

C/C++

```
# pragma omp parallel
{
  #pragma omp for
  for (...) {...}
}
```

```
# pragma omp parallel for
for(){...}
```

* Since Fortran do has a block structure “!\$omp end do” Is not required

Work-Sharing: Sections

```
1 #pragma omp parallel sections
2 {
3     #pragma omp section
4     {
5         work_1();
6     }
7     #pragma omp section
8     { work_2(); }
9 }
```

Line 1 Team of threads formed (parallel region).

Line 3-8 One thread is working on each section.

Line 9 End of parallel sections with an implied barrier.

Scales only to the number of sections.

Work-Sharing: Sections

```
1 !$omp parallel sections
2
3   !$omp section
4
5     call work_1()
6
7   !$omp section
8     call work_2()
9
10 !$omp end parallel sections
```

Line 1 Team of threads formed (parallel region).

Line 3-9 One thread is working on each section.

Line 10 End of parallel sections with an implied barrier.

Scales only to the number of sections.

Shorthand (again)

```
1 !$OMP PARALLEL DO  
2     do i=1,N  
3         a(i) = b(i) + c(i)  
4     enddo  
5 !$OMP END PARALLEL DO
```

```
1 !$OMP PARALLEL  
2 !$OMP DO  
3     do i=1,N  
4         a(i) = b(i) + c(i)  
5     enddo  
6 !$OMP END DO  
7 !$OMP END PARALLEL
```

! Create team of threads
! Assign iterations to threads

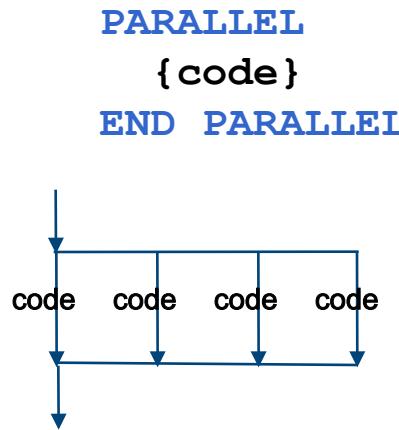
But, how does all this work
if we need other code/work here?

OpenMP Parallel Constructs

Replicated : Work blocks are executed by all threads.

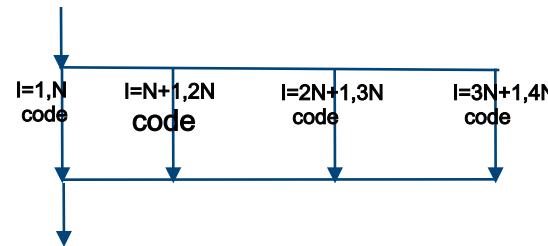
Work-Sharing : Work is divided among threads.

4 Threads



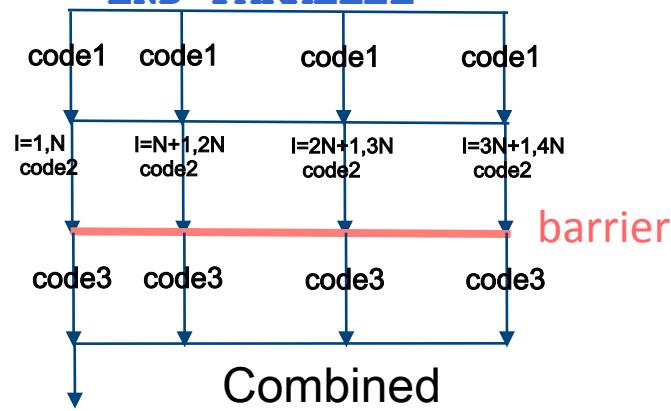
Replicated

```
PARALLEL DO
do I = 1,N*4
{code}
end do
END PARALLEL DO
```



Work-Sharing

```
PARALLEL
{code1}
DO
do I = 1,N*4
{code2}
end do
END DO
{code3}
```



Combined

OpenMP Clauses

Clauses control the behavior of an OpenMP directive:

1. Schedule for for/do worksharing (Static, Dynamic, Guided, etc.)
2. Data scoping (Private, Shared, Default)
3. Initialization (e.g. COPYIN, FIRSTPRIVATE)
4. Parallelize a region or not (IF(logic_expression))
5. Number of threads used (NUM_THREADS)

Schedule Clause for loop worksharing

schedule(static)

Each CPU receives one set of contiguous iterations

```
!$omp do schedule(...)  
do i=1,128  
  A(i)=B(i)+C(i)  
enddo
```

schedule(static, C)

Iterations are divided round-robin fashion in chunks of size C

schedule(dynamic, C)

Iterations handed out in chunks of size C as CPUs become available

schedule(guided, C)

Each of the iterations are handed out in pieces of exponentially decreasing size, with C minimum number of iterations to dispatch each time

schedule (runtime)

Schedule and chunk size taken from the OMP_SCHEDULE environment variable

Example - schedule(static,16), threads = 4

```
!$omp parallel do schedule(static,16)
do i=1,128
  A(i)=B(i)+C(i)
enddo
```

<u>thread0</u> :	do i=1,16 A(i)=B(i)+C(i) enddo do i=65,80 A(i)=B(i)+C(i) enddo
<u>thread1</u> :	do i=17,32 A(i)=B(i)+C(i) enddo do i = 81,96 A(i)=B(i)+C(i) enddo
<u>thread2</u> :	do i=33,48 A(i)=B(i)+C(i) enddo do i = 97,112 A(i)=B(i)+C(i) enddo
<u>thread3</u> :	do i=49,64 A(i)=B(i)+C(i) enddo do i = 113,128 A(i)=B(i)+C(i) enddo

Comparison of Scheduling Options

name	type	chunk	chunk size	chunk #	static or dynamic	compute overhead
partitioned	static	no	N/P	P	static	lowest
interleaved	static	yes	C	N/C	static	low
simple dynamic	dynamic	optional	C	N/C	dynamic	medium
guided	guided	optional	decreasing from N/P*	no fewer than C	dynamic	high
runtime	runtime	no	varies	varies	varies	varies

*Decreases as
(N-unassigned)/P

OpenMP Data Environment

- Clauses control the data-sharing attributes of variables within a parallel region:

shared, private, reduction (firstprivate, lastprivate)

Default variable scope (in parallel region):

1. Variables declared in main/program (C/F90) are shared by default
2. Global variables are shared by default
3. Automatic variables within subroutines called from within a parallel region are private (reside on a stack private to each thread)
4. Loop index of worksharing loops are private.
5. Default scoping rule can be changed with **default** clause

Private & Shared Data

shared - Variable is shared (seen) by all threads

private - Each thread has a private instance (copy) of the variable

Defaults: The for-loop index is private, all other variables are shared

```
!$omp parallel do shared(a,b,c,n)    private(i)      OK to be explicit;  
                                         but not necessary.  
                                         do i = 1,n{  
                                         a(i) = b(i) + c(i)  
                                         enddo
```

All threads have access to the same storage areas for a, b, c, and n, but each loop has its own private copy of the loop index, i

Private & Shared Data

shared - Variable is shared (seen) by all threads

private - Each thread has a private instance (copy) of the variable

Defaults: The for-loop index is private, all other variables are shared

```
#pragma omp parallel for shared(a,b,c,n)  private(i)
    for (i=0; i<n; i++){
        a[i] = b[i] + c[i];
    }
```

OK to be explicit;
but not necessary.

All threads have access to the same storage areas for a, b, c, and n, but each loop has its own private copy of the loop index, i

Private Data Example

- In the following loop, each thread needs its own private copy of temp
- If temp were shared, the result would be unpredictable since each thread would be writing/reading to/from the same memory location

```
#pragma omp parallel for shared(a,b,c,n) private(temp,i)
for (i=0; i<n; i++){
    temp = a[i] / b[i];
    c[i] = temp + cos(temp);
}
```

- A **lastprivate(temp)** clause will copy the last loop(stack) value of temp to the (global) temp storage when the parallel for is complete.
- A **firstprivate(temp)** would copy the global temp value to each stack's temp.

Private Data Example

- In the following loop, each thread needs its own private copy of temp
- If temp were shared, the result would be unpredictable since each thread would be writing/reading to/from the same memory location

```
!$omp parallel for shared(a,b,c,n) private(temp,i)
do i = 1,n
    temp = a(i) / b(i)
    c(i) = temp + cos(temp)
endo
```

- A **lastprivate(temp)** clause will copy the last loop(stack) value of temp to the (global) temp storage when the parallel DO is complete.
- A **firstprivate(temp)** would copy the global temp value to each stack's temp.

- Operation that combines multiple elements to form a single result
- A variable that accumulates the result is called a reduction variable
- In parallel loops reduction operators and variables must be declared

```
float asum=0.0, aprod=1.0;

#pragma omp parallel for reduction(:asum) reduction(*:aprod)
for (i=0; i<n; i++){
    asum = asum + a[i];
    aprod = aprod * a[i];
}
```

Each thread has a private **asum** and **aprod**, initialized to the operator's identity

- After the loop execution, the master thread collects the private values of each thread and finishes the (global) reduction

Reduction

- Operation that combines multiple elements to form a single result
- A variable that accumulates the result is called a reduction variable
- In parallel loops reduction operators and variables must be declared

```
real asum=0.0, aprod=1.0

!$omp parallel do reduction(:asum) reduction(*:aprod)
do i = 1,n
    asum = asum + a(i)
    aprod = aprod * a(i)
enddo
print*, asum, aprod
```

Each thread has a private **asum** and **aprod**, initialized to the operator's identity

- After the loop execution, the master thread collects the private values of each thread and finishes the (global) reduction

Synchronization

- Synchronization is used to impose order constraints and to protect access to shared data
- High-Level Synchronization
 - critical
 - atomic
 - barrier
 - ordered
- Low-Level Synchronization
 - locks

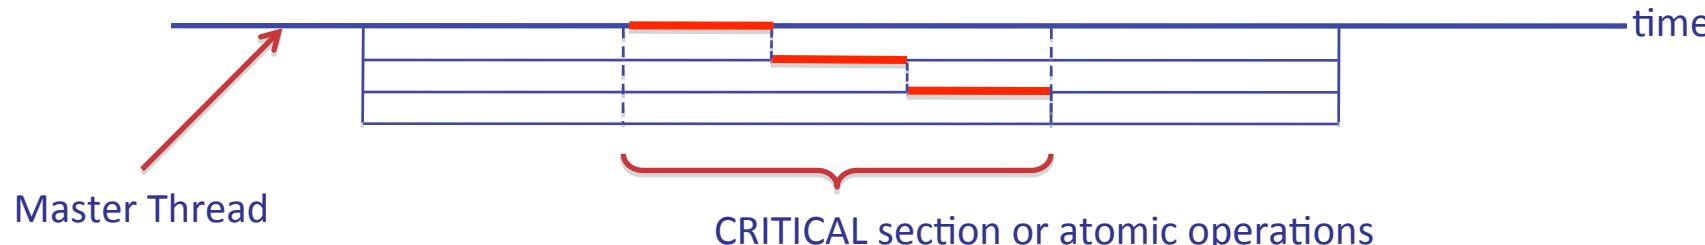
Synchronization: Critical/Atomic Directives

- When each thread must execute a section of code serially the region must be marked with **critical/end critical** directives
- Use the **#pragma omp atomic** directive for simple cases: can use hardware support

```
#pragma omp parallel shared(sum,x,y)
{
    ...
    #pragma omp critical
        update(x);
        update(y);
        sum=sum+1;
    }
    ...
}
```

```
#pragma omp parallel shared(sum)
{
    ...
    #pragma omp atomic
        sum=sum+1
    ...
}
```

Atomic has
read,
write,
update,
capture
clauses.



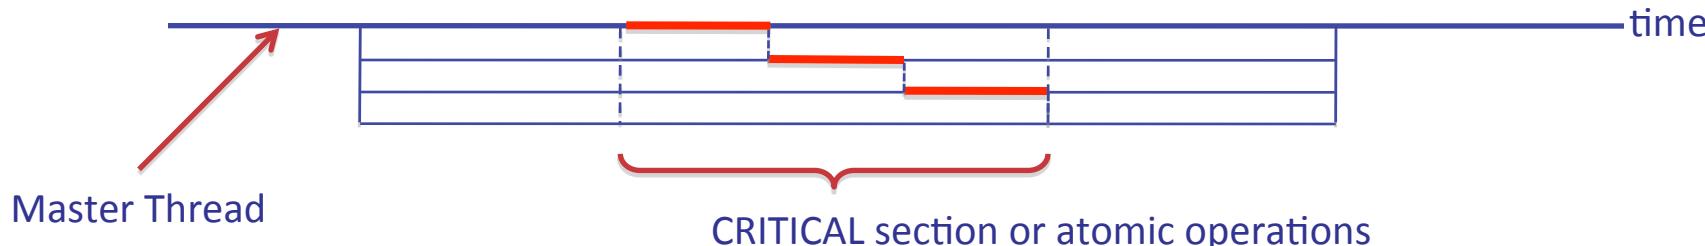
Synchronization: Critical/Atomic Directives

- When each thread must execute a section of code serially the region must be marked with **critical/end critical** directives
- Use the **!\$omp atomic** directive for simple cases: can use hardware support

```
!$omp parallel shared(sum,x,y)
...
 !$omp critical
 update(x);
 update(y);
 sum=sum+1;
 !$omp end critical
 ...
 !$omp end parallel
```

```
!$omp parallel shared(sum)
...
 !$omp atomic
 sum=sum+1;
 ...
 !$omp end parallel
```

Atomic has
read,
write,
update,
capture
clauses.



Synchronization: Barrier

- Barrier: Each thread waits until all threads arrive and a flush occurs

```
#pragma omp parallel shared (A, B, C) private(id)
{
    id=omp_get_thread_num();
    A[id] = big_calc1(id);
    #pragma omp barrier
    #pragma omp for
    for(i=0;i<N;i++) {
        C[i]=big_calc3(i,A);
    }←Implicit barrier
    #pragma omp for nowait
    for(i=0;i<N;i++) {
        B[i]=big_calc2(C, i);
    }←No implicit barrier due to nowait
    A[id] = big_calc4(id);
}←Implicit barrier
```

Synchronization: Barrier

- Barrier: Each thread waits until all threads arrive and flush occurs

```
!$omp parallel shared (A, B, C) private(id)

    id=omp_get_thread_num()
    A(id) = big_calc1(id)
    !$omp barrier

    !$omp do
        do i = 1,N; C(i)=big_calc3(i,A); enddo
    !$omp end do <----- Implicit barrier

    !$omp do
        do i = 1,N; B(i)=big_calc2(C, i); enddo
    !$omp end do nowait <----- No implicit barrier due to nowait
    A(id) = big_calc4(id);
    !$omp end parallel <----- Implicit barrier
```

Barriers

- The previous example could have been done with multiple parallel regions:
 1. There is an implicit barrier at the end of a parallel region.
 2. However, creating and recreating thread teams is expensive.

Synchronization: Ordered

- The ordered region executes in the sequential order

```
#pragma omp parallel private (tmp)
#pragma omp for ordered reduction(+:countVal)
for (i=0;i<N;i++) {
    tmp = foo(i);
    #pragma omp ordered
    countVal+= consume(tmp);
}
```

Synchronization: Ordered

- The ordered region executes in the sequential order

```
!$omp parallel private (tmp)
!$omp do ordered reduction(:countVal)
do i =1,N{
    tmp = foo(i)
    !$omp ordered
    countVal = countVal + consume(tmp)
end do
```

NOWAIT

- When a work-sharing region is exited, a barrier is implied - all threads must reach the barrier before any can proceed.
- By using the NOWAIT clause at the end of each loop inside the parallel region, an unnecessary synchronization of threads can be avoided.

```
#pragma omp parallel
{
    #pragma omp for nowait
    {
        for (i=0; i<n; i++)
            {work(i);}
    }
    #pragma omp for schedule(guided,k)
    {
        for (i=0; i<m; i++)
            {x[i]=y[i]+z[i];}
    }
}
```

NOWAIT

- When a work-sharing region is exited, a barrier is implied - all threads must reach the barrier before any can proceed.
- By using the NOWAIT clause at the end of each loop inside the parallel region, an unnecessary synchronization of threads can be avoided.

```
!$OMP PARALLEL
  !$OMP DO
    do i=1,n
      work(i)
    enddo
  !$OMP END DO NOWAIT
  !$OMP DO schedule(guided,k)
    do i=1,m
      x(i)=y(i)+z(i)
    enddo
  !$OMP END DO
 !$OMP END PARALLEL
```

Runtime Library Routines

function	description
omp_get_num_threads()	Number of threads in team, N
omp_get_thread_num()	Thread ID {0 -> N-1}
omp_get_num_procs()	Number of machine CPUs
omp_in_parallel()	True if in parallel region & multiple thread executing
omp_set_num_threads(#)	Set the number of threads in the team

Environment Variables

variable	description
OMP_NUM_THREADS = <i>integer</i>	Set to default no. of threads to use
OMP_SCHEDULE =“ <i>schedule-type</i> [, <i>chunk_size</i>]”	Sets “runtime” in loop schedule clause: “...omp for/do schedule(runtime) ”
OMP_DISPLAY_ENV = <i>anyvalue</i>	Prints runtime environment at beginning of code execution.

[...] = optional

47

OpenMP Wallclock Timers

```
real*8 :: omp_get_wtime,    omp_get_wtick()          (Fortran)
double    omp_get_wtime(),   omp_get_wtick();           (C)
```

```
double t0, t1, dt, res;
...
t0 = omp_get_wtime();
<work>
t1 = omp_get_wtime();
dt = t1 - t0;

res = 1.0/omp_get_wtick();
printf("Elapsed time = %lf\n",dt);
printf("clock resolution = %lf\n",res);
```

NUM_THREADS clause

- Use the **NUM_THREADS** clause to specify the number of threads to execute a parallel region

```
#pragma omp parallel num_threads(scalar int expression)
{
    <code block>
}
```

where **scalar integer expression** must evaluate to a positive integer

- `num_threads()` supersedes the number of threads specified by the **OMP_NUM_THREADS** environment variable or that set by the **OMP_SET_NUM_THREADS** function

NUM_THREADS clause

- Use the **NUM_THREADS** clause to specify the number of threads to execute a parallel region

```
!$omp parallel num_threads(scalar integer expression)
    <code block>
!$omp end parallel
```

where **scalar integer expression** must evaluate to a positive integer

- `num_threads()` supersedes the number of threads specified by the **OMP_NUM_THREADS** environment variable or that set by the **OMP_SET_NUM_THREADS** function

Mutual Exclusion: Lock Routines

When each thread must execute a section of code serially, **locks** provide a **more flexible** way of ensuring serial access than **CRITICAL** and **ATOMIC** directives

```
call OMP_INIT_LOCK(maxlock)
!$OMP PARALLEL SHARED(x,y)
...
call OMP_set_lock(maxlock)
call update(x)
call OMP_unset_lock(maxlock)
...
!$OMP END PARALLEL
call OMP_DESTROY_LOCK(maxlock)
```

OpenMP 3.0

- First update to the spec since 2005
- Tasking: move beyond loops with generalized tasks and support complex and dynamic control flows
- Loop collapse: combine nested loops automatically to expose more concurrency
- Enhanced loop schedules: Support aggressive compiler optimizations of loop schedules and give programmers better runtime control over the kind of schedule used
- Nested parallelism support: better definition of and control over nested parallel regions, and new API routines to determine nesting structure

Tasks Parallelism

- Allows to parallelize irregular problems
- OpenMP had tasks internally: chunks of execution; with OpenMP
- Tasks depend on other tasks: dataflow
- You can declare many tasks, OpenMP executes them, satisfying dependencies

Task dependencies

```
1 ! In parallel single region here. Single thread generates tasks
! Other threads of team execute tasks.
2 for (i=0; i<children; i++)
3 #pragma omp task
4   f(i)
5 #pragma omp taskwait
6 parent_function()
```

Children tasks are put in a queue
can be executed in any order
but before the parent

Tasks: Usage

Task Construct:

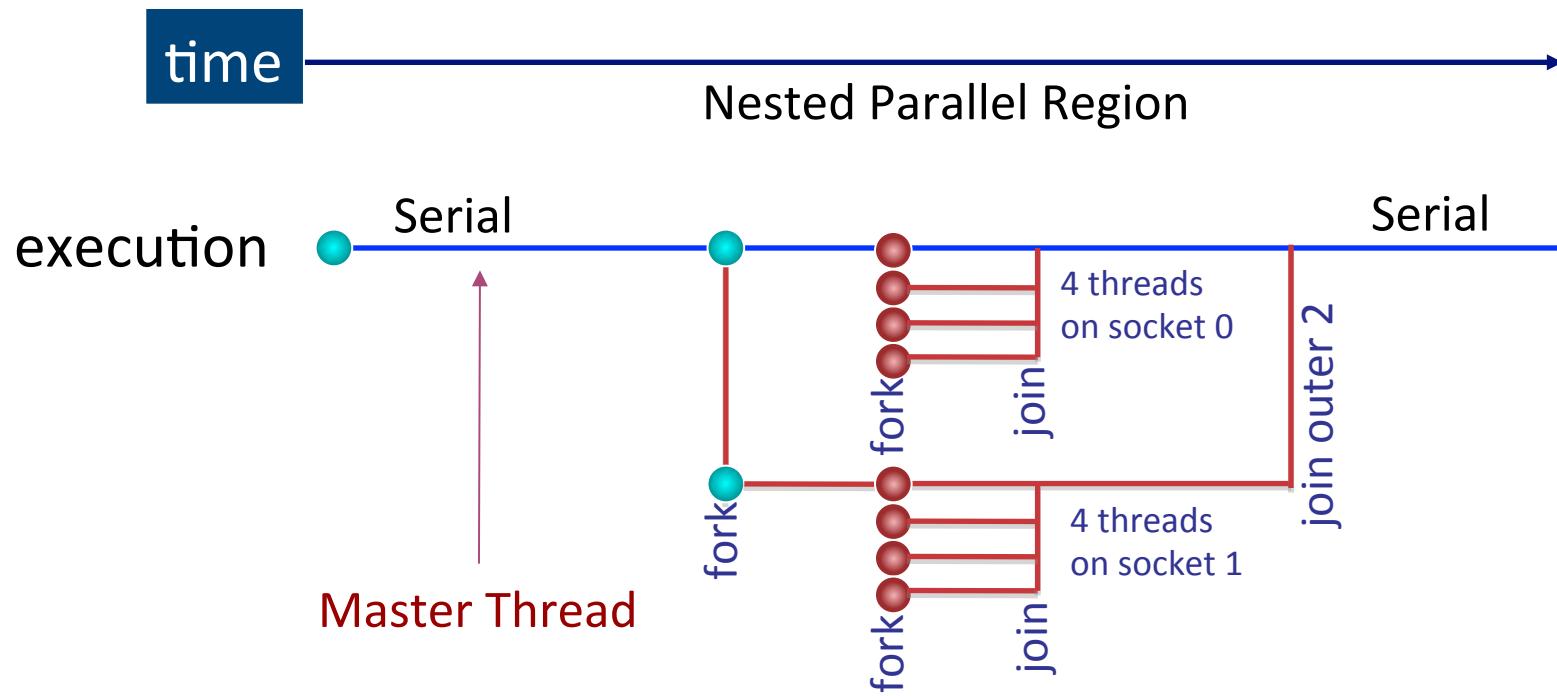
```
#pragma omp task [clause[,]clause] ...]
```

structured-block

where clause can be

- Data scoping clauses
 - **shared (list), private (list), firstprivate (list), default(shared | none)**
- Scheduling clauses
 - **untied**
- Other clauses
 - **if (expression)**

Loop Nesting



While OpenMP 3.0 supports nested parallelism, many implementations may ignore the nesting by serializing the inner parallel regions

Loop Collapse

- Allow collapsing of perfectly nested loops
- Will form a single loop and then parallelize it:

```
#pragma omp parallel do collapse(2)
for(i=0;i<n;i++){
    for(j=0;j<n;j++) {
        ....
    }
}
```

Loop Collapse

- Allow collapsing of perfectly nested loops
- Will form a single loop and then parallelize it:

```
!$omp parallel do collapse(2)
do i=1,n
    do j=1,n
        .....
    end do
end do
```

References

- <http://www.openmp.org/>
- *Parallel Programming in OpenMP*, by Chandra, Dagum, Kohr, Maydan, McDonald, Menon
- *Using OpenMP*, by Chapman, Jost, Van der Pas (OpenMP2.5)
- http://www.nic.uoregon.edu/iwomp2005/iwomp2005_tutorial_openmp_rvdp.pdf
- <http://webct.ncsa.uiuc.edu:8900/public/OPENMP/>

FOR FORTRAN USERS and backups

Parallel Regions & Modes

There are two OpenMP “modes”

- **static** mode
 - Fixed number of threads -- set in the **OMP_NUM_THREADS** env.
Or the threads may be set by a function call (or clause) inside the code:
 - **omp_set_num_threads** runtime function
num_threads(#) clause
- **dynamic** mode:
 - Number of threads can change under OS control from one parallel region to another using:

Note: the user can only define the maximum number of threads, compiler can use a smaller number

Default variable scoping (Fortran example)

```
Program Main
  Integer, Parameter :: nmax=100
  Integer :: n, j
  Real*8 :: x(n,n)
  Common /vars/ y(nmax)
  ...
  n=nmax; y=0.0
  !$OMP Parallel do
    do j=1,n
      call Adder(x,n,j)
    end do
  ...
End Program Main
```

```
Subroutine Adder(a,m,col)
  Common /vars/ y(nmax)
  SAVE array_sum
  Integer :: i, m
  Real*8 :: a(m,m)

  do i=1,m
    y(col)=y(col)+a(i,col)
  end do
  array_sum=array_sum+y(col)

End Subroutine Adder
```

Default data scoping in Fortran (cont.)

Variable	Scope	Is use safe?	Reason for scope
n	shared	yes	declared outside parallel construct
j	private	yes	parallel loop index variable
x	shared	yes	declared outside parallel construct
y	shared	yes	common block
i	private	yes	parallel loop index variable
m	shared	yes	actual variable <i>n</i> is shared
a	shared	yes	actual variable <i>x</i> is shared
col	private	yes	actual variable <i>j</i> is private
array_sum	shared	no	declared with SAVE attribute

Kent Milfeld

Lars Koesterke

Antonio Gomez

milfeld/lars/agomez@tacc.utexas.edu

For more information:
www.tacc.utexas.edu

