



OpenMP Tasking Concepts

TACC OpenMP Team

milfeld/lars/agomez@tacc.utexas.edu

Sides at: tinyurl.com/tacc-openmp or portal.tacc.utexas.edu/training click View Details



THE UNIVERSITY OF TEXAS AT AUSTIN
TEXAS ADVANCED COMPUTING CENTER

Learning Objective

-- Tasking --

- Review Parallel and Worksharing Concepts
 - Forking, data scoping -- for/do construct, implicit barriers
- Basic Task Syntax and Operations
- Task Synchronization
- Running Tasks in Parallel
- Data-sharing and firstprivate Default for Tasks
- Common Use Cases for Tasks
- Task Dependences
- Taskloop

OpenMP Pre 3.0

Data Parallel Paradigms

- About mapping **threads** to loop-chunks of work
- About work independence for **worksharing**
 - Determine iteration independence through Read/Write (RW) analysis: mainly
RaW, WaR, (WaW of concern & RaR not of concern)
- About avoiding **race conditions**.

RaW = Read after Write, etc. 3

Review Parallel and Worksharing Concepts

Forking, data scoping -- for/do construct, implicit barriers

Basic Task Syntax and Operations

Task Synchronization and Operations

Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop

Parallel & Worksharing

- **parallel** Construct
 - Forks Threads
 - creates (forks) threads, sets data environment
 - Implied barrier at the end of a region.
- **do/for, single, and sections** Constructs
 - **workshares** iterations/block(s) of code
 - do/for
 - Uses prescribed method (static, **dynamic**, etc.) to schedule threads.
 - **iteration limit required**

Worksharing

- Splits work up and gives **CHUNKS OF INDEPENDENT** work to a team of threads.
- Threads wait at an **IMPLIED BARRIER**.
- **DATA ENVIRONMENT** is set by parallel region & altered by workshare constructs: sections/single/do/for. (If you have a team of threads, you must be in a parallel region.)

Limitation of Worksharing

- Worksharing **requires** ability to know number of work units before execution (**loop count**).
- **Dynamic scheduling** of parallel loop work-chunks **is limited**. Think of work chunks being queued to execute from a FIFO queue*.
- Ideal for “data parallel”, but **“task parallel” requires inconvenient construction**.

*This can be somewhat limited, but reasonable for monotonically increasing/decreasing work per iteration.

Review Parallel and Worksharing Concepts

Forking, data scoping -- for/do construct, implicit barriers

Basic Task Syntax and Operations

Task Synchronization

Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop

Getting Started with Tasking

- Easy to begin using for novice
- Uncomplicated syntax and specification for many cases
- Task clauses provide enough flexibility to be efficient (in programming syntax and execution) for more complicated use cases.

Is it easy? – that depends.
It is somewhat like worksharing—
There are simple and complex cases.

What is Tasking for?

- Irregular Computing:
 - Dynamic Execution (better control of work chunks)
 - Nested Execution
 - Can Employ Dependences (& Priorities)
- Cases for Irregular Computing:
 - Execute independent iterations of **while loop** in parallel
 - **Follow pointers** until a NULL pointer is reached, performing independent work at each pointer position.
 - Note: pointer chasing is inherently serial but work at each pointer position must be independent.
 - Follow pointers & perform independent work at nodes in a **graph tree**
 - **Order executions** that have task (work) **dependences**

Generating a task

- We first go over generating tasks in a serial region – computationally impractical.
- But...it is a practical base for developing concepts.
- Hold on– we'll get to executing tasks concurrently!

Generating a task

- **task** is a directive, with clauses (not shown)
 - Encountering thread generates a task (the block of work); by default the task is “deferrable”. It can execute the task immediately or defer it.
 - Deferred tasks are queued to be executed after generation, allowing generating task to continue.

Creating two tasks in a serial region.

Tasks do independent work.

C/C++

```
#pragma omp task
    foo(j);

#pragma omp task
    for(i=0;i<n;i++){...};
```

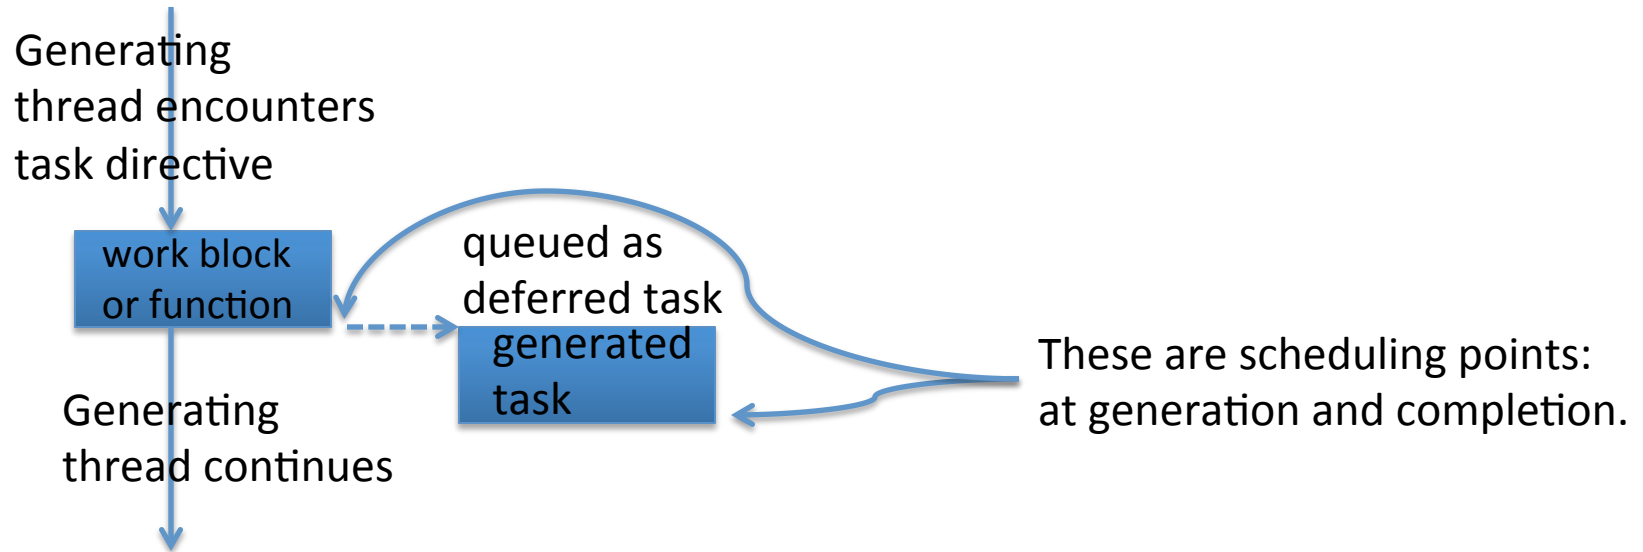
F90

```
!$omp task
    foo(j)
!$omp end task

!$omp task
    do i = 1,n; ... ;enddo
!$omp end task
```

Deferred Task


Deferred task.



This thread (at a thread scheduling point) or another thread can execute the queued task.

Immediate Task

Immediate task.



Tasking Syntax & Operation Explained

- If the previous code is in a serial section of program (usually not the case for a task), the task will probably execute immediately –i.e. will not run the generated task later as a deferred task (asynchronously).
- “Deferrable” means that it can be executed later (by throwing it on the queue)—this is not required though; it may execute the task immediately (by the task generating thread) if it sees this as optimal.
- Hence, because there is just one thread available, the runtime may execute the task immediately after generation. Also, it may be submitted to the queue as a deferred task. In this latter case it will most likely be pulled from the queue and executed right away, because the runtime realizes there are no other threads to execute it.

What is a Task

- Work generated by a **task construct**– yeah, yeah we know that now! This IS an **explicit task**.
- Threads of a **parallel** region, executing the replicated work-- these are also tasks, with the distinctive name: **implicit tasks***.
- Standard Spec: a task is, a “specific instance of executable code and its data environment” generated by”...

In presentation TASK will only mean Explicit Task unless otherwise stated.

*“Each implicit task is tied to a different thread in the team” – more on this later.

16

Review Parallel and Worksharing Concepts

Forking, data scoping -- for/do construct, implicit barriers

Basic Task Syntax and Operations

Task Synchronization

Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop

Synchronizing tasks (sibling tasks)

- If a task is deferred, use the **taskwait** construct to wait for completion at some point in the code.

C/C++

```
#pragma omp task
{ foo(j); }

#pragma omp task
{ for(i=0;i<n;i++){...}; }

#pragma omp taskwait
```

F90

```
!$omp task
    call foo(j)
!$omp end task

!$omp task
    do i = 1,n; ... ;enddo
!$omp end task
!$omp taskwait
```

Synchronizing tasks (nested tasks)

- A **taskgroup** construct waits for all sibling and their descendants. (Yes, tasks can generate child tasks.)

C/C++

```
#pragma taskgroup
{
    #pragma omp task
    foo(j);

    #pragma omp task
    { for(i=0;i<n;i++)
        #pragma omp task
        foo(i);
    }
}
```

F90

```
!omp taskgroup
!$omp task
    call foo(j)
!$omp end task

!$omp task
    do i = 1,n
        !$omp task
            call foo(i);
        !$omp end task
    enddo
!$omp end task
!$omp end taskgroup
```

Nested
Tasks

19

Review Parallel and Worksharing Concepts

Forking, data scoping -- for/do construct, implicit barriers

Basic Task Syntax

Task Synchronization

Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop

Generating Concurrent Tasks— Running in a parallel region

- First, create a team of threads, to work on tasks.
- Use a single thread to generate tasks.

C/C++

```
#pragma omp parallel
{
    #pragma omp single
    {
        //generate multiple tasks
    }
}
```

F90

```
!$omp parallel

    !$omp single

        !generate multiple tasks


    !omp end single
!$omp end parallel
```

Generate multiple tasks: with loop (while/do while/do/for) or recursion

21

Tasks in parallel region

- Threads of Team will dequeue & execute tasks
- Shared variables of parallel region are also shared by tasks
- Tasks obey explicit and implicit barriers, also taskwait & taskgroup



Review Parallel and Worksharing Concepts

Forking, data scoping -- for/do construct, implicit barriers

Basic Task Syntax and Operations

Task Synchronization

Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop

Tasks: Basic Data-sharing Attributes

- If the task generating construct is in a parallel region any shared variables remain shared.
- for/do construct variables of a worksharing loop are private (see spec.)
- Private variables is on an enclosing construct are firstprivate for the task.

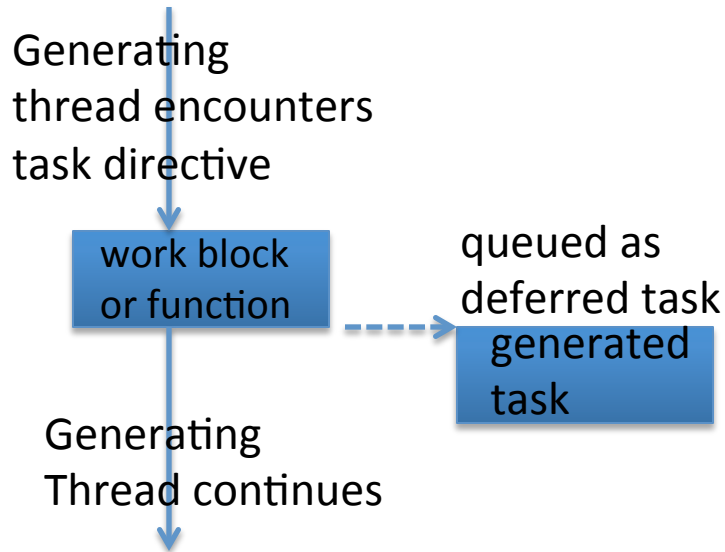
Corner Cases:

Otherwise, if a variable data-sharing attribute is not determined by prescribed rules, it is firstprivate.

24

Deferred Task

Deferred task.



Basic Concept:

Any thread can pick up a task– but argument/var values at the generation time are needed to determine the work AFTER task is dequeued.

In a parallel region j is shared, and hence needs to be declared `firstprivate`.

Corner Case: In a serial region j IS NOT shared and hence will be automatically `firstprivate`.

```
#pragma omp task
foo(j);

j++;
```

```
!$omp task
foo(j)
!$omp end task

j = j+1
```

Scheduling Optimization

- For a small number of threads and tasks, and a large diversity in task work—an imbalance will occur. Even with moderate diversity and large thread and task counts, an imbalance may be present. The **priority clause** can alleviate this problem.

```
for (i=0;i<N; i++) {  
    #pragma omp task priority(i)  
    compute_array(array[i], M);  
}
```

```
do i=1,N  
    !$omp task priority(i)  
    call compute_array(matrix(:, i), N)  
    !$omp end task  
enddo
```

26

Review Parallel and Worksharing Concepts

Forking, data scoping -- for/do construct, implicit barriers

Basic Task Syntax and Operations

Task Synchronization

Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop

While loop

- firstprivate clause is necessary, since *cntr* is shared and value must be captured for work.

```
int cntr = 100;
#pragma omp parallel
#pragma omp single

while(cntr>0) {
    #pragma omp task \
        firstprivate(cntr)

    printf("cntr=%d\n",cntr);

    cntr--;
}
```

```
integer cntr = 100
!$omp parallel
!$omp single

do while(cntr>0)
    !$omp task &
        firstprivate(cntr)

        print*, "cntr= ",cntr
    !$omp end task

    cntr = cntr - 1
enddo
!$omp end single
!$omp end parallel
```

28

Exploiting tasks within while loop

- The loop is executed **SERIALLY**, but concurrently with the dequeued tasks.
 - So, the non-tasking loop parts should not be costly.
 - Any generated tasks can be picked up directly by other team members.

```
#pragma omp single
while(cnt>0){
    #pragma omp task firstprivate(cnt)
    printf("cnt=%d\n",cnt);

    cnt--;
}
```



Serial – generation
parallel – Tasking
Serial – incrementer

Pointer Chasing

- ***ptr*** points to a C/C++ structure or F90 defined type

```
int *ptr;  
...//initialize pointer  
#pragma omp parallel  
#pragma omp single  
  
while(ptr) {  
    #pragma omp task firstprivate(ptr)  
    process(ptr) ;  
  
    ptr = ptr->next;  
}
```

```
integer,pointer :: ptr  
...! initialize pointer  
!$omp parallel  
!$omp single  
  
do while(associated(ptr))  
    !$omp task firstprivate(ptr)  
    process(ptr)  
    !$omp end task  
  
    ptr = ptr%next  
enddo  
  
!$omp end single  
!$omp end parallel
```

30

Using Recursion in C/C++

- Starting point for tree searches.
- Note, the recursion function, not the “process”, is directly tasked.

```
#pragma omp parallel
#pragma omp single
    chase(ptr) ;

...
void chase(node *ptr) {

    if(ptr->next) {
        #pramga omp task
        chase(ptr->next) ;
    }
    process(ptr)
}
```

Just another way
to build a queue
of tasks— but with
recursion (bunch
of context switching)

Undeferred Tasks with **if** clause

```
while(ptr) {  
  
    usec=ptr->cost*factor;  
  
    #pramga omp task if(usec>0.01) firstprivate(ptr)  
    process(ptr)  
  
    ptr = ptr->next;  
}
```

If execution time (usec) is less than 0.01 microseconds, the **if** argument is false:

- Generating thread will suspend generation
- Generating thread will execute the task
- Generating thread will resume generation

Tied Tasks (default)

- Tasks can be stopped and continued (at scheduling points). Use `taskyield` directive to create your own scheduling point.
- A Tied Task can only be continued by the thread that initially began the execution of the task. This can be important for keeping cached data available (hot) to the thread, and for locks.

Untied Tasks

Syntax: `... omp task untied`

- Any thread may continue execution of the task
- Important for controlling task generation when queue becomes full:
A single untied generator task can let another thread continue.

1. When queue is full, the task generation thread may execute queued tasks.
2. If it gets stuck, another thread can execute the untied generation task.

Untied Task Example

```
#pragma omp task untied
{
    for(i=0; i<Nlarge;i++)
        #pragma omp task
        process(i);
}
```

Details: read this later:

- A single (outer) generating task generates a large number of tasks. These are put on a queue to execute.
- When the queue is filled, the runtime will ask the thread executing the task generation to switch at the next task generation (task constructs are task switching points) and execute a generated task (a queued task). That is, suspend the task generating task and pick up tasks from the queue.
- If there is an imbalance in work and the thread that was executing task generation gets stuck in processing one of the generated tasks, and other threads many finish their tasks. Because the “stuck” thread is tied to the generating task, no other thread can pick up the generation task to continue putting tasks on the queue.
- However, if the (outer) generating task is untied, any other thread can resume the generation task and continue to generate.

35

Review Parallel and Worksharing Concepts

Forking, data scoping -- for/do construct, implicit barriers

Basic Task Syntax and Operations

Task Synchronization

Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop

Depend clause

- Dependences are derived from *dependence-type* and the *list* items of the depend clause.

Syntax: ... task depend(*dependence-type*: *list*)

dependence-type:

in If a storage location of at least one list item is the same as a list item appearing with an **out** or **inout** *dependence-type* of a previously generated task, the **in** task will be dependent on the **out/inout** task.

out
inout If a storage location of at least one list item is the same as a list item appearing with an **in**, **out** or **inout** *dependence-type* of a previously generated task, then the generated task will be a dependent task of the previously generated task.

Comment: What?

Depend clause (in other words)

- When a task requires a previously computed object, it has an **input (in) dependence**.
- When a task is computing a value for a object that may be required later, it has an **output (out)dependence**.
- When a tasking requires a previous computed object and computes a new value it has an **input-output (inout) dependence**.
- Dependences are formed through the **object list** – **variable, element references, array sections**, etc.
The list items may be objects used in the routine.

Comment: OK, this is starting to make sense.

38

Depend clause (another approach – RW)

- Enforces scheduling constraints of sibling tasks (or ordered loop iterations – not discussed here).


Syntax: ... task depend(*dependence-type*: *list*)

dependence-type: == what the task needs to do

in	think of as a read
out	think of as a write
inout	think of as a read and then a write

<i>list</i>	Item it needs to “Read” or “Write”
-------------	------------------------------------

- Dependences



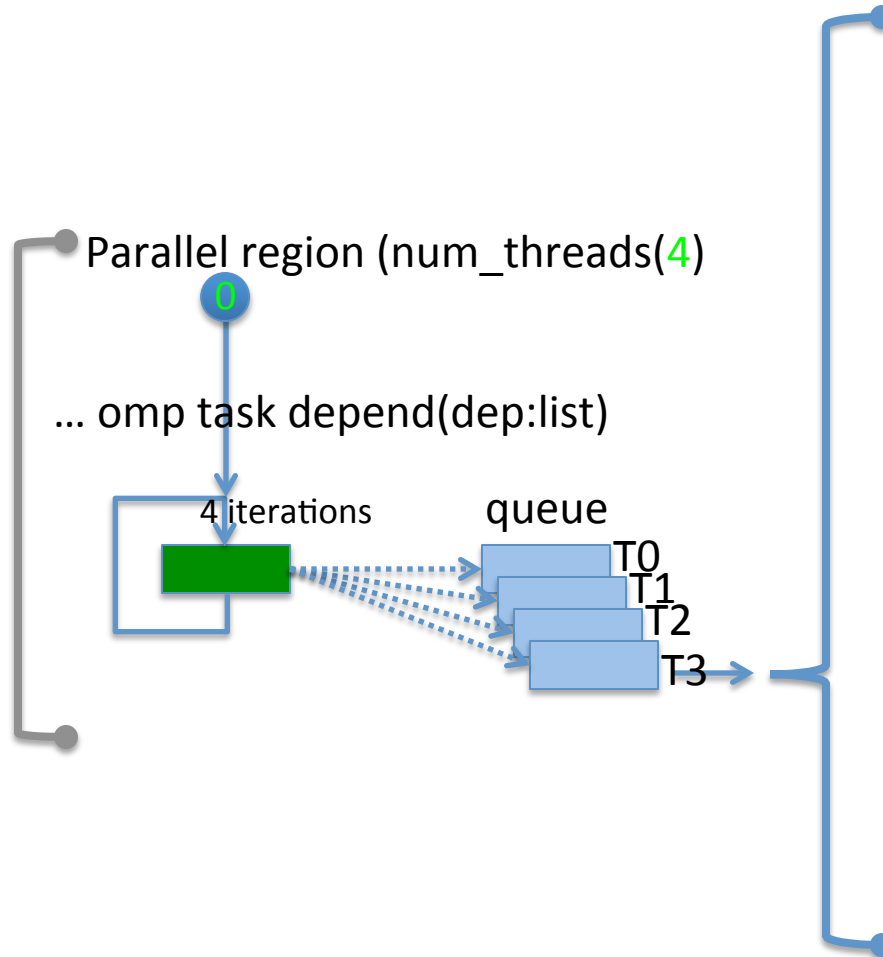
Task 1 execution dependence

When runtime executes T0, it checks previously generated tasks for identical list times.

There are no previous tasks— so this task has NO dependence. EVEN if it has an IN (read) dependence type!

- Dependences

Task 3 execution dependence



T3 checks previously generated tasks for identical list times.

If an identical list item exists in previously generated tasks, T3 adheres to the *dependence-type*.

Task depend clause

- Flow Control (RaW, Read after Write)

```
x = 1;
#pragma omp parallel
#pragma omp single
{
    #pragma omp task shared(x) depend(out: x)
    x = 2;
    #pragma omp task shared(x) depend(in: x)
    printf("x = %d\n", x);
}
...
```

T1 is put on queue,
sees no previously
queued tasks
with x identifier →
No dependences.

T2 is put on queue,
sees previously queued
task with identifier →
Has RaW dependence.

Print value is
always 2.

Task depend clause

- Anti-dependence (WaR, write after read)

```
x = 1;
#pragma omp parallel
#pragma omp single
{
    #pragma omp task shared(x) depend( in: x)
    printf("x = %d\n", x);
    #pragma omp task shared(x) depend(out: x)
    x = 2;
}
...
```

T1 is put on queue,
sees no previously
queued tasks
with x identifier →
No dependences.

T2 is put on queue,
sees previously queued
task with identifier →
has WaR dependence.

Print value is
always 1.

Task depend clause

- Output Dependence (WaW, Write after Write)

```
x = 1;
#pragma omp parallel
#pragma omp single
{
    #pragma omp task shared(x) depend(out: x)
    printf("x = %d\n", x);
    #pragma omp task shared(x) depend(out: x)
    x = 2;
}
```

T1 is put on queue,
sees no previously
queued tasks
with x identifier →
No dependences.

T2 is put on queue,
sees previously queued
task with identifier →
has WaW dependence.

Print value is
always 1.

Task depend clause

- (RaR, no dependence)

```
x = 1;
#pragma omp parallel
#pragma omp single
{
    #pragma omp task shared(x) depend(in: x)
    printf("x = %d\n", x);
    #pragma omp task shared(x) depend(in: x)
    x = 2;
}
```

T1 is put on queue,
sees no previously
queued tasks
with x identifier →
No dependences.


T2 is put on queue,
sees previously queued
task with x identifier →
has NO ordering
(because it is RAR)

Print value is
1 or 2.

Task depend clause

- Following a graph


```
#pragma omp parallel
#pragma omp single
{
T1  #pragma omp task depend(out:a,b,c)
    f1(&a, &b, &c);
T2  #pragma omp task depend(out:d,e)
    f2(&d, &e);
T3  #pragma omp task depend(in:c,d)
    f3(c,d);
T4  #pragma omp task depend(out,e)
    f4(&e);
T5  #pragma omp task depend(in:a) depend(out:b,c)
    f5(a,&b,&c)
}
```



Task Depend Clause

- Following non-computed variables-- works, too.

```
#pragma omp parallel
#pragma omp single
{
T1 #pragma omp task depend(out:t1,t2,t3)
    f1(&a, &b, &c);
T2 #pragma omp task depend(out:t4,t5)
    f2(&d, &e);
T3 #pragma omp task depend(in:t3,t4)
    f3(c,d);
T4 #pragma omp task depend(out,t5)
    f4(&e);
T5 #pragma omp task depend(in:t1)depend(out:t2,t3)
    f5(a,&b,&c)
}
```



Review Parallel and Worksharing Concepts

Forking, data scoping -- for/do construct, implicit barriers

Basic Task Syntax and Operations

Task Synchronization

Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop

taskloop

- Iterations of loops are executed as tasks (of a taskgroup)
 - Single generator needed
 - All team members are not necessary
 - Implied taskgroup for the construct

Syntax: `... omp taskloop [clauses]`

Some clauses:

grainsize		number of iterations assigned to a task.
numtasks	or	create this number of tasks
	default:	number of tasks & iterations/task implementation defined
untied		tasks need not be continued by initial thread of task
nogroups		don't create a task group
priority		for each task (default 0)

Taskloop

- ptr points to a structure in C/C++, defined type in F90

```
void parallel_work(void) { // execute by single in parallel
    int i, j;
    #pragma omp taskgroup
    {
        #pragma omp task
        long_running(); // can execute concurrently

        #pragma omp taskloop private(j) grainsize(500) nogroup
        for (i = 0; i < 10000; i++) //can execute concurrently
            for (j = 0; j < i; j++)
                loop_body(i, j);
    } // end taskgroup
}
```

Summary

- Tasks are used mainly in irregular computing.
- Tasks are often generated by a single thread.
- Task generation can be recursive.
- Depend clause can prescribe dependence.
- Priority provides hint for execution order.
- First private is default data-sharing attribute, shared variable remain shared.
- Untied generator task can assure generation progress.